

# Building Neural Networks with scikit-learn

---

INTRODUCING NEURAL NETWORKS IN SCIKIT-LEARN



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**scikit-learn support for neural networks**

**scikit-learn vs. deep learning frameworks**

**Perceptrons and neurons**

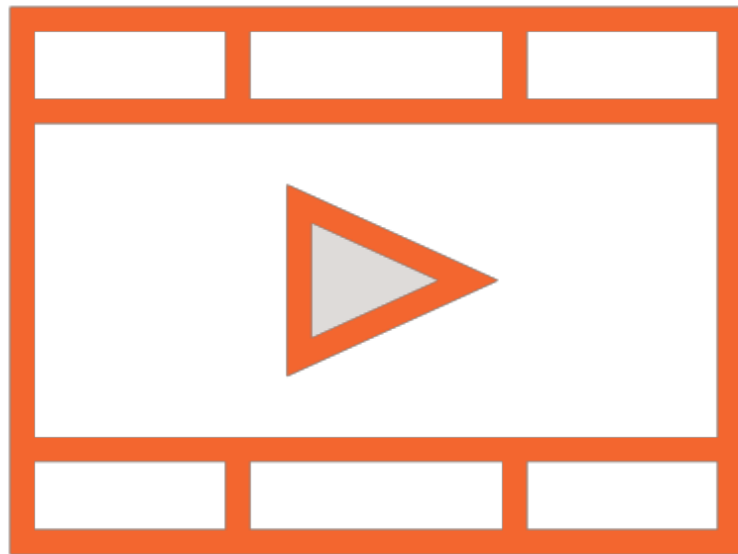
**Multi-layer perceptrons (MLPs) and neural networks**

**Training a neural network**

# Prerequisites and Course Outline

---

# Prerequisites




**Basic Python programming**

**Prior ML exposure, basic ML workflow**

**Building and training ML models in  
scikit-learn**

# Prerequisite Courses



**Building Your First scikit-learn Solution**

**Building Regression Models with scikit-learn**

**Building Classification Models with scikit-learn**

# Course Outline



**Neural networks in scikit-learn**

**Regression and classification with neural networks**

**Text and image classification**

**Dimensionality reduction with Restricted Boltzmann Machines**

# Neural Networks in scikit-learn

---

Support for neural networks in  
scikit-learn is currently quite  
limited



# scikit-learn vs. Other Frameworks

## scikit-learn

Most popular library for general purpose ML

Vast array of estimators for classification, regression, clustering

Implemented using **traditional ML algorithms**

Very **limited support** for building neural networks

## TensorFlow, PyTorch

Widely used libraries that specialize in deep learning

Relatively small number of algorithmic estimators for those problems

Very little support for traditional ML algorithms

Entirely focused on building neural networks

# Neural Networks on scikit-learn

## scikit-learn

Work directly with Pandas data frames  
and NumPy arrays

No specialized GPU support

Not suited to distributed training

## TensorFlow, PyTorch

Work with special data types called  
**tensors** for multidimensional arrays

Help leverage power of GPUs

Extensive support for distributed training

# Neural Networks on scikit-learn

## scikit-learn

Limited number of neural network building blocks

Support for just **fully-connected neural networks with regularization**

Impossible to build complex RNNs, CNNs

No pre-trained models for transfer learning

## TensorFlow, PyTorch

Large numbers of neuron types, activation functions, loss functions

Building blocks to support different kinds of neural network layers

Relatively simple to build complex RNN and CNN architectures

Impressive array of pre-trained models available for transfer learning

# Neural Networks in scikit-learn

**Supervised**

**Unsupervised**

# Neural Networks in scikit-learn

**Multi-layer Perceptrons  
(MLP)**

**Restricted Boltzmann  
Machines (RBM)**

# Perceptrons and Neurons

---

# Neural Networks in scikit-learn

**Supervised**

**Unsupervised**

# Neural Networks in scikit-learn

**Multi-layer Perceptrons  
(MLP)**

**Restricted Boltzmann  
Machines (RBM)**



# Neural Networks in scikit-learn

**Multi-layer Perceptrons  
(MLP)**

**Restricted Boltzmann  
Machines (RBM)**

# Perceptron



**Simplest Artificial Neural Network architecture**

**Originally invented in 1957 by Frank Rosenblatt**

**Precursor** to the neuron used today

# Perceptron



**Calculates the weighted sum of inputs**

**Applies a step function with a threshold**

**Output - positive**

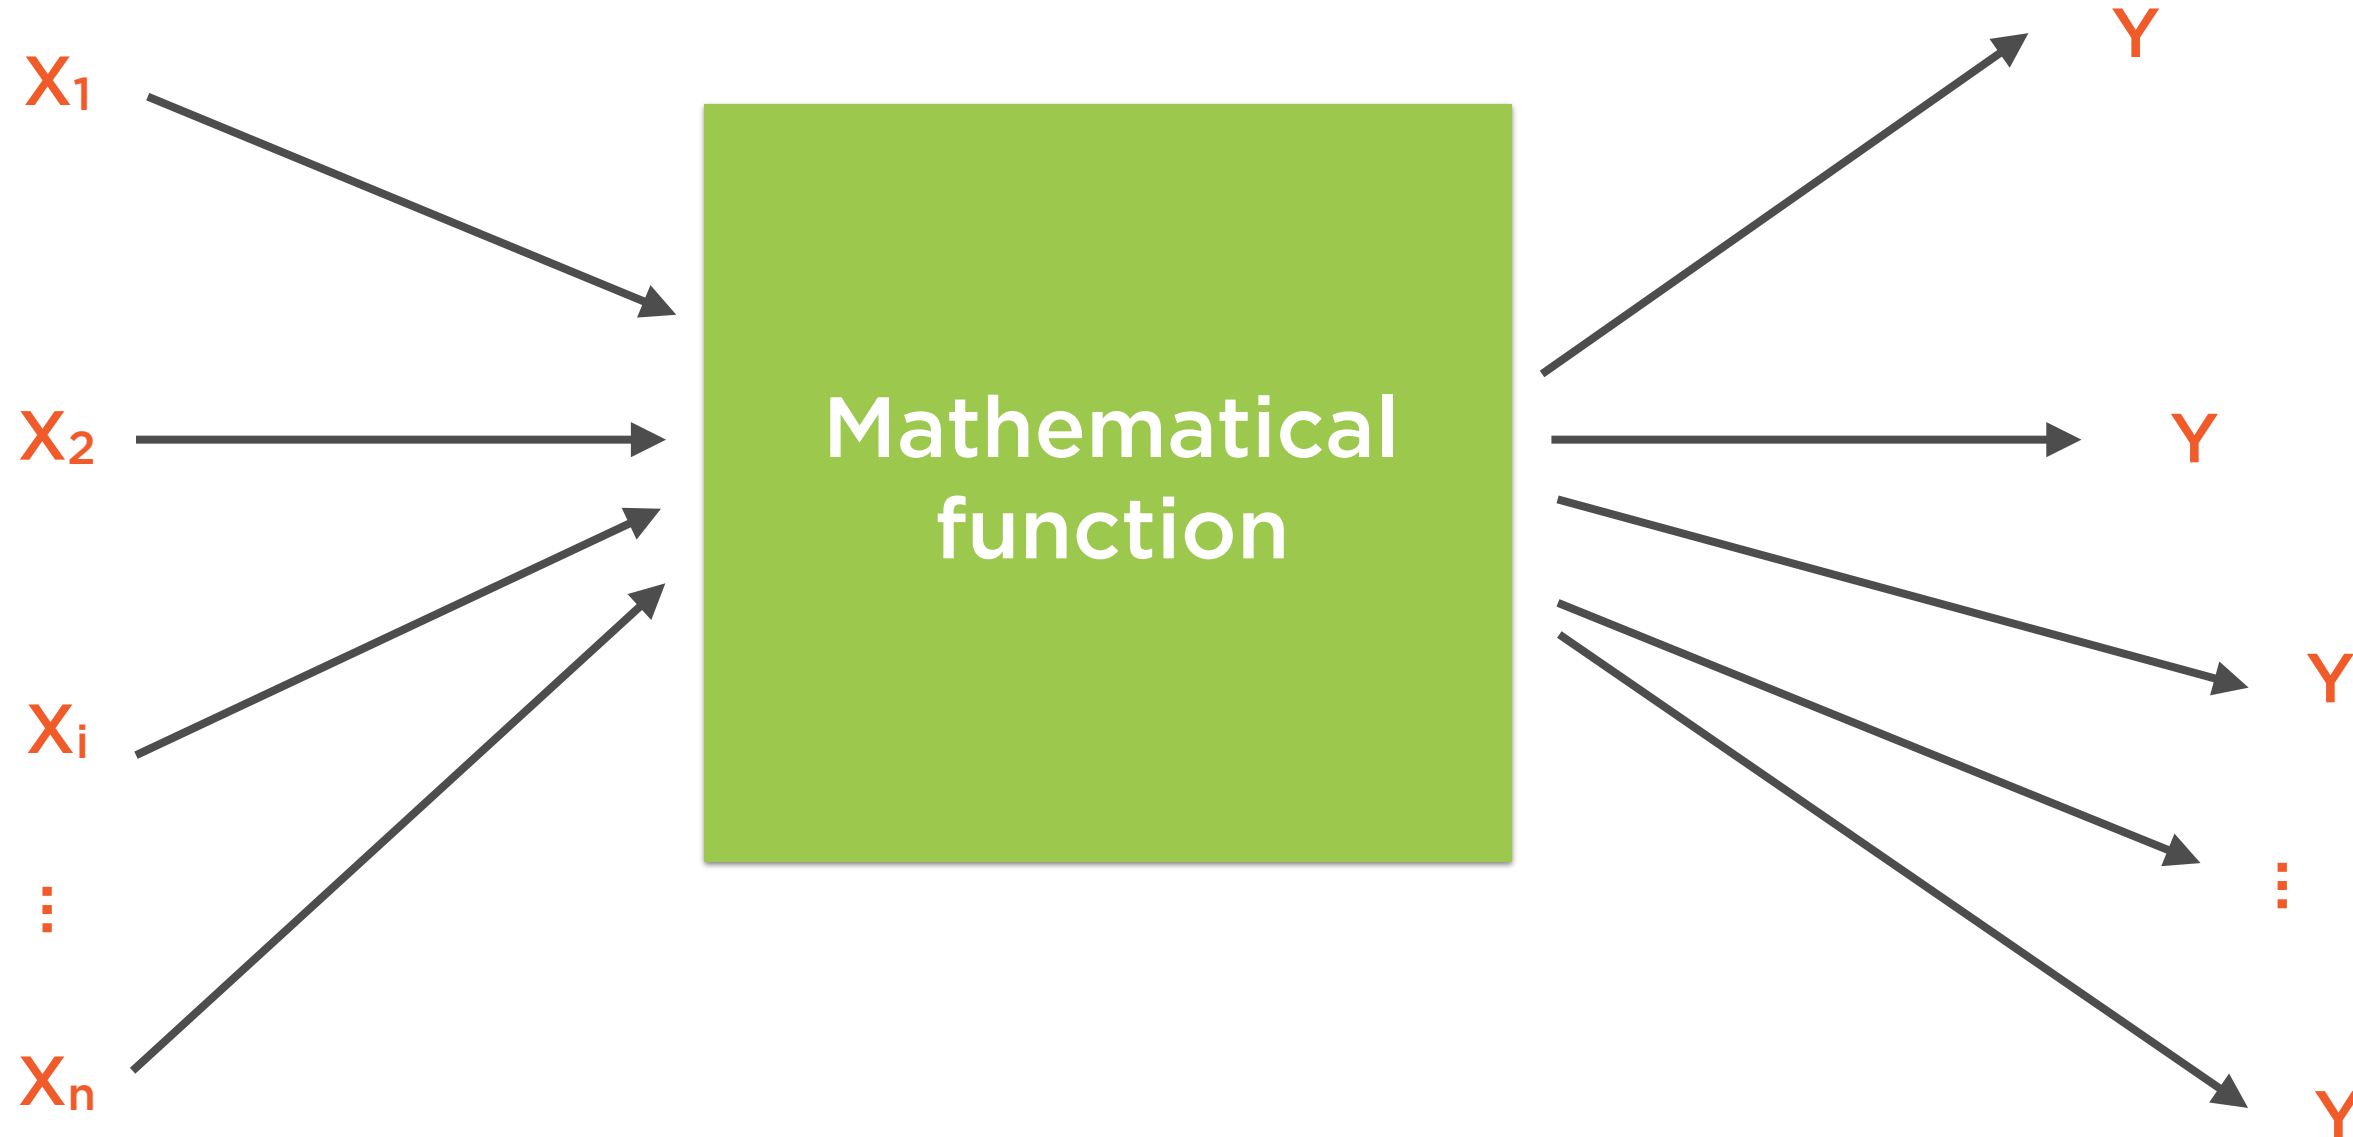
- if value **above** threshold

**Output - negative**

- if value **below** threshold


Perceptron ~ Neuron with a step  
activation function

# Operation of a Single Neuron




**For an active neuron a change in inputs should trigger a corresponding change in the outputs**

# Operation of a Single Neuron




**The outputs of neurons feed into the neurons from the next layer**

# Operation of a Single Neuron



**Each connection is associated with a weight**

# Operation of a Single Neuron

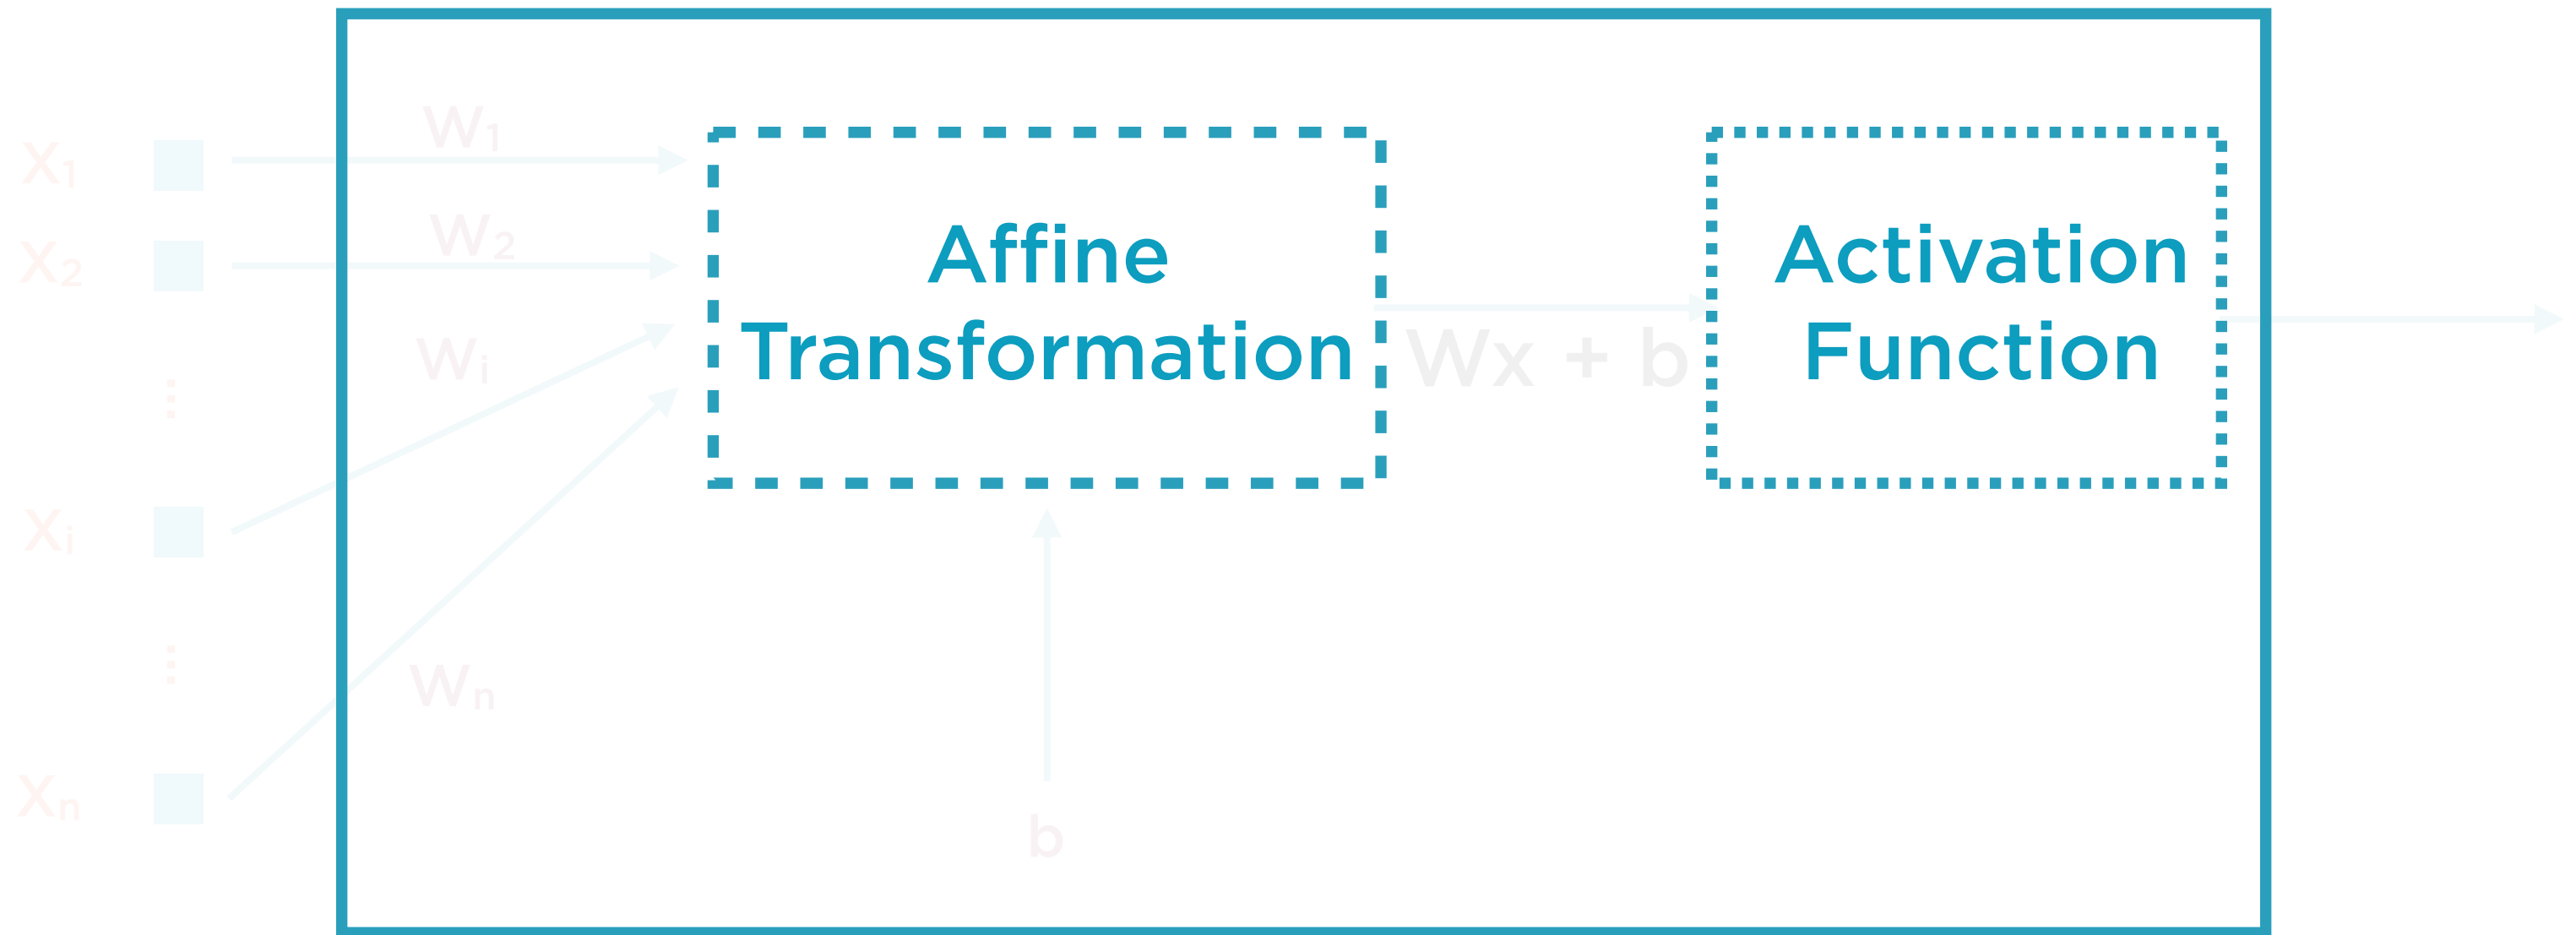


If the second neuron is sensitive to the output of the first neuron, the **connection between them gets stronger**

**$W$  increases**

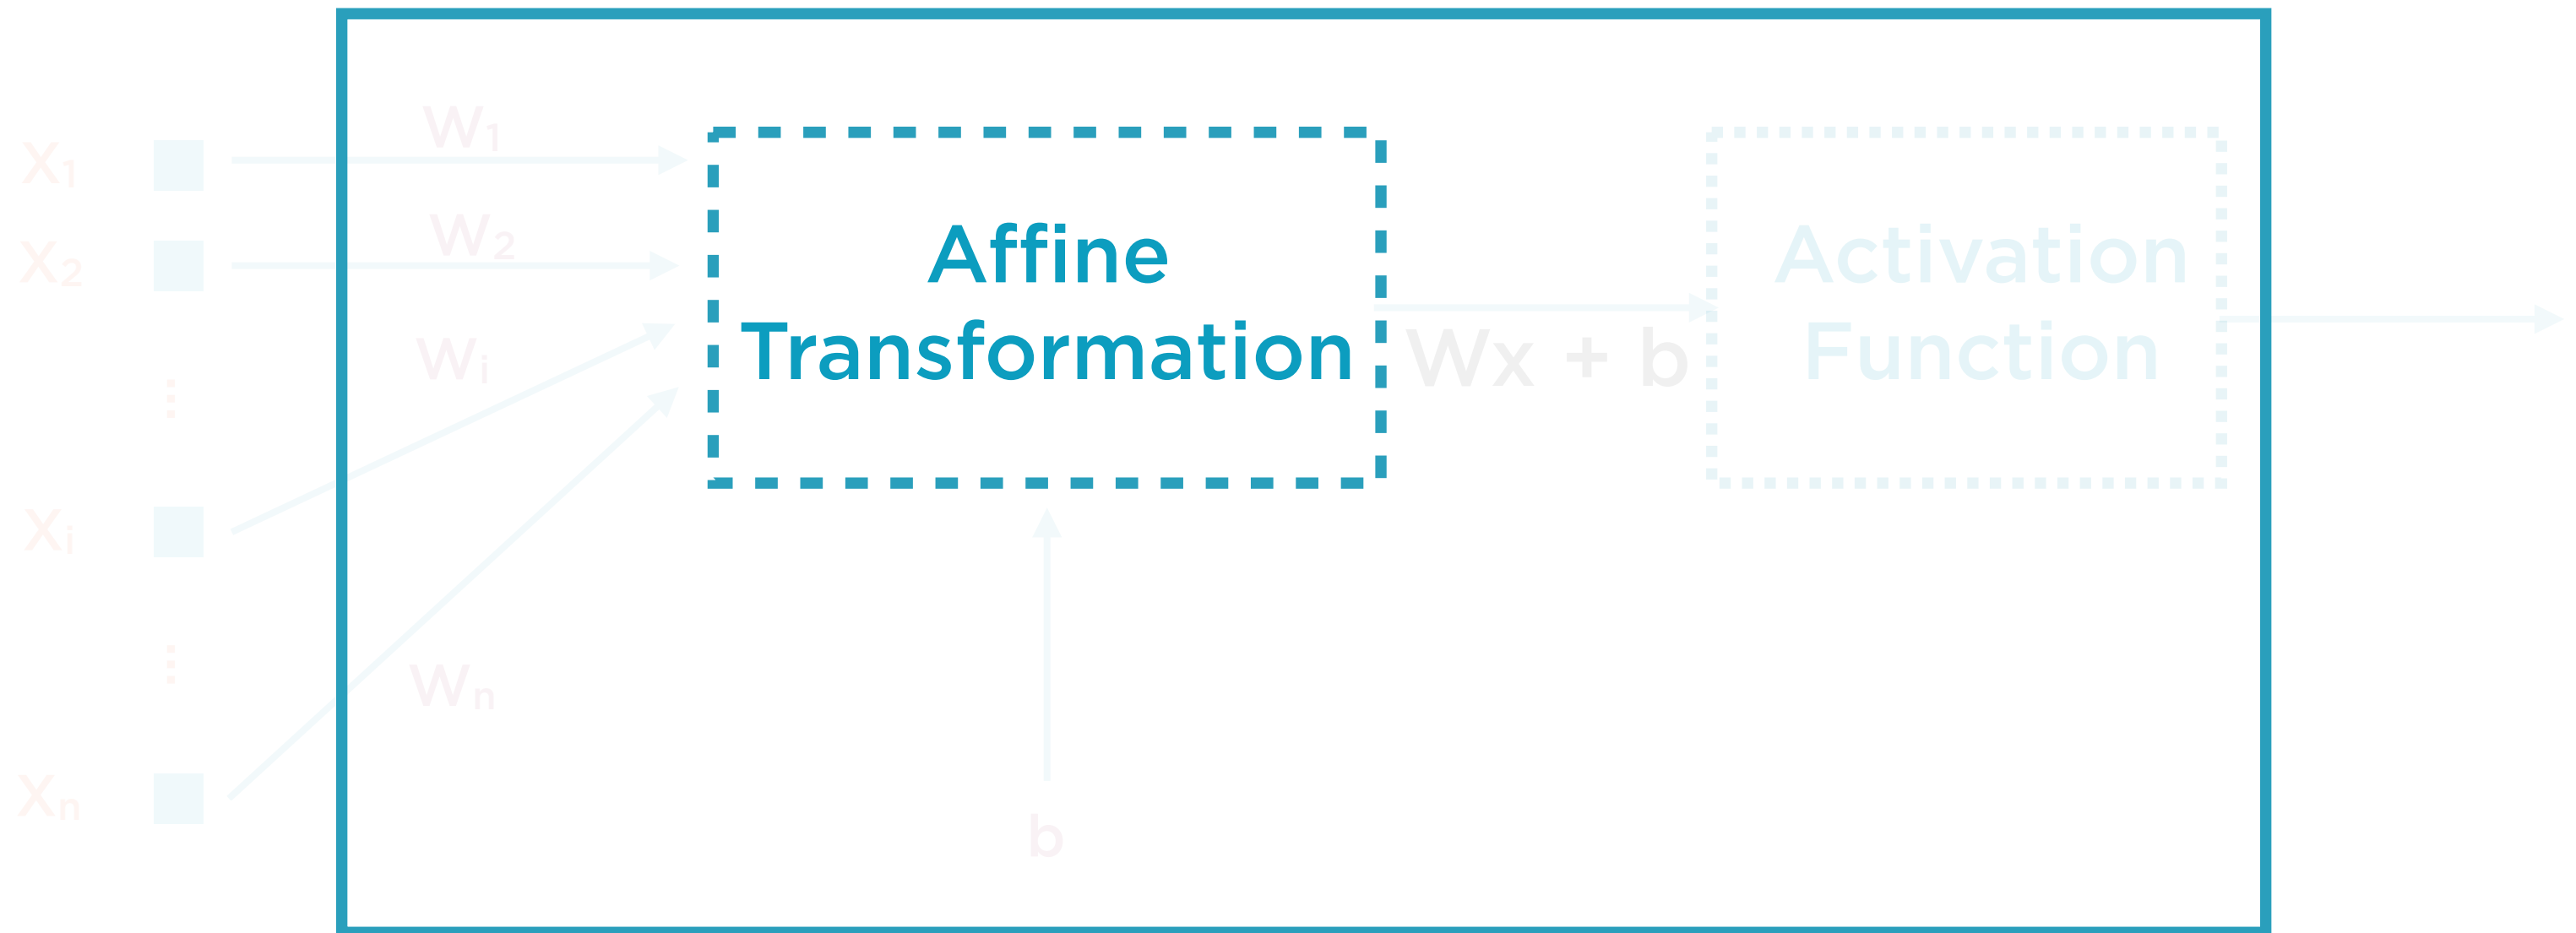


# Operation of a Single Neuron



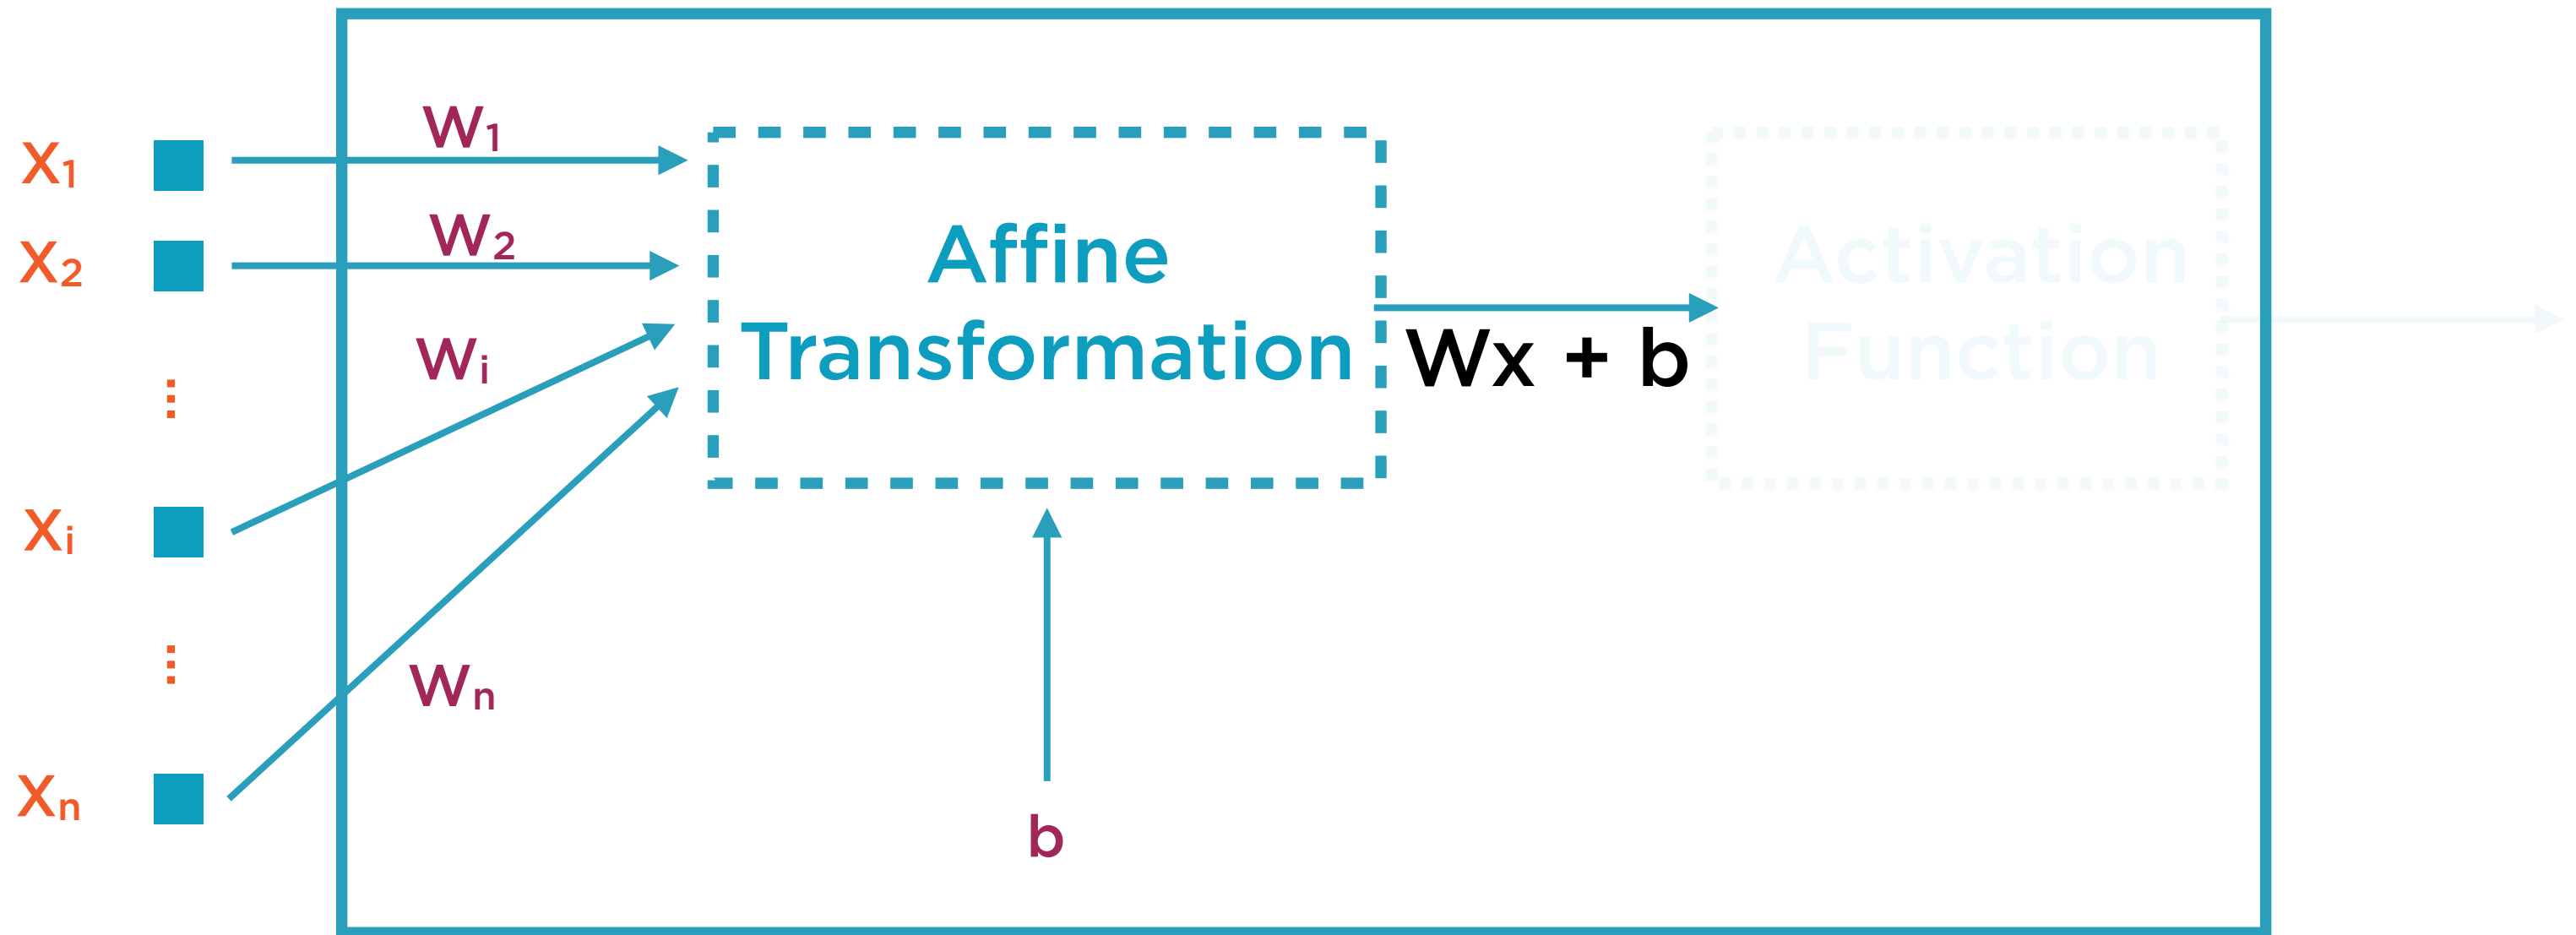
**Each neuron only applies two simple functions to its inputs**

# Affine Transformation



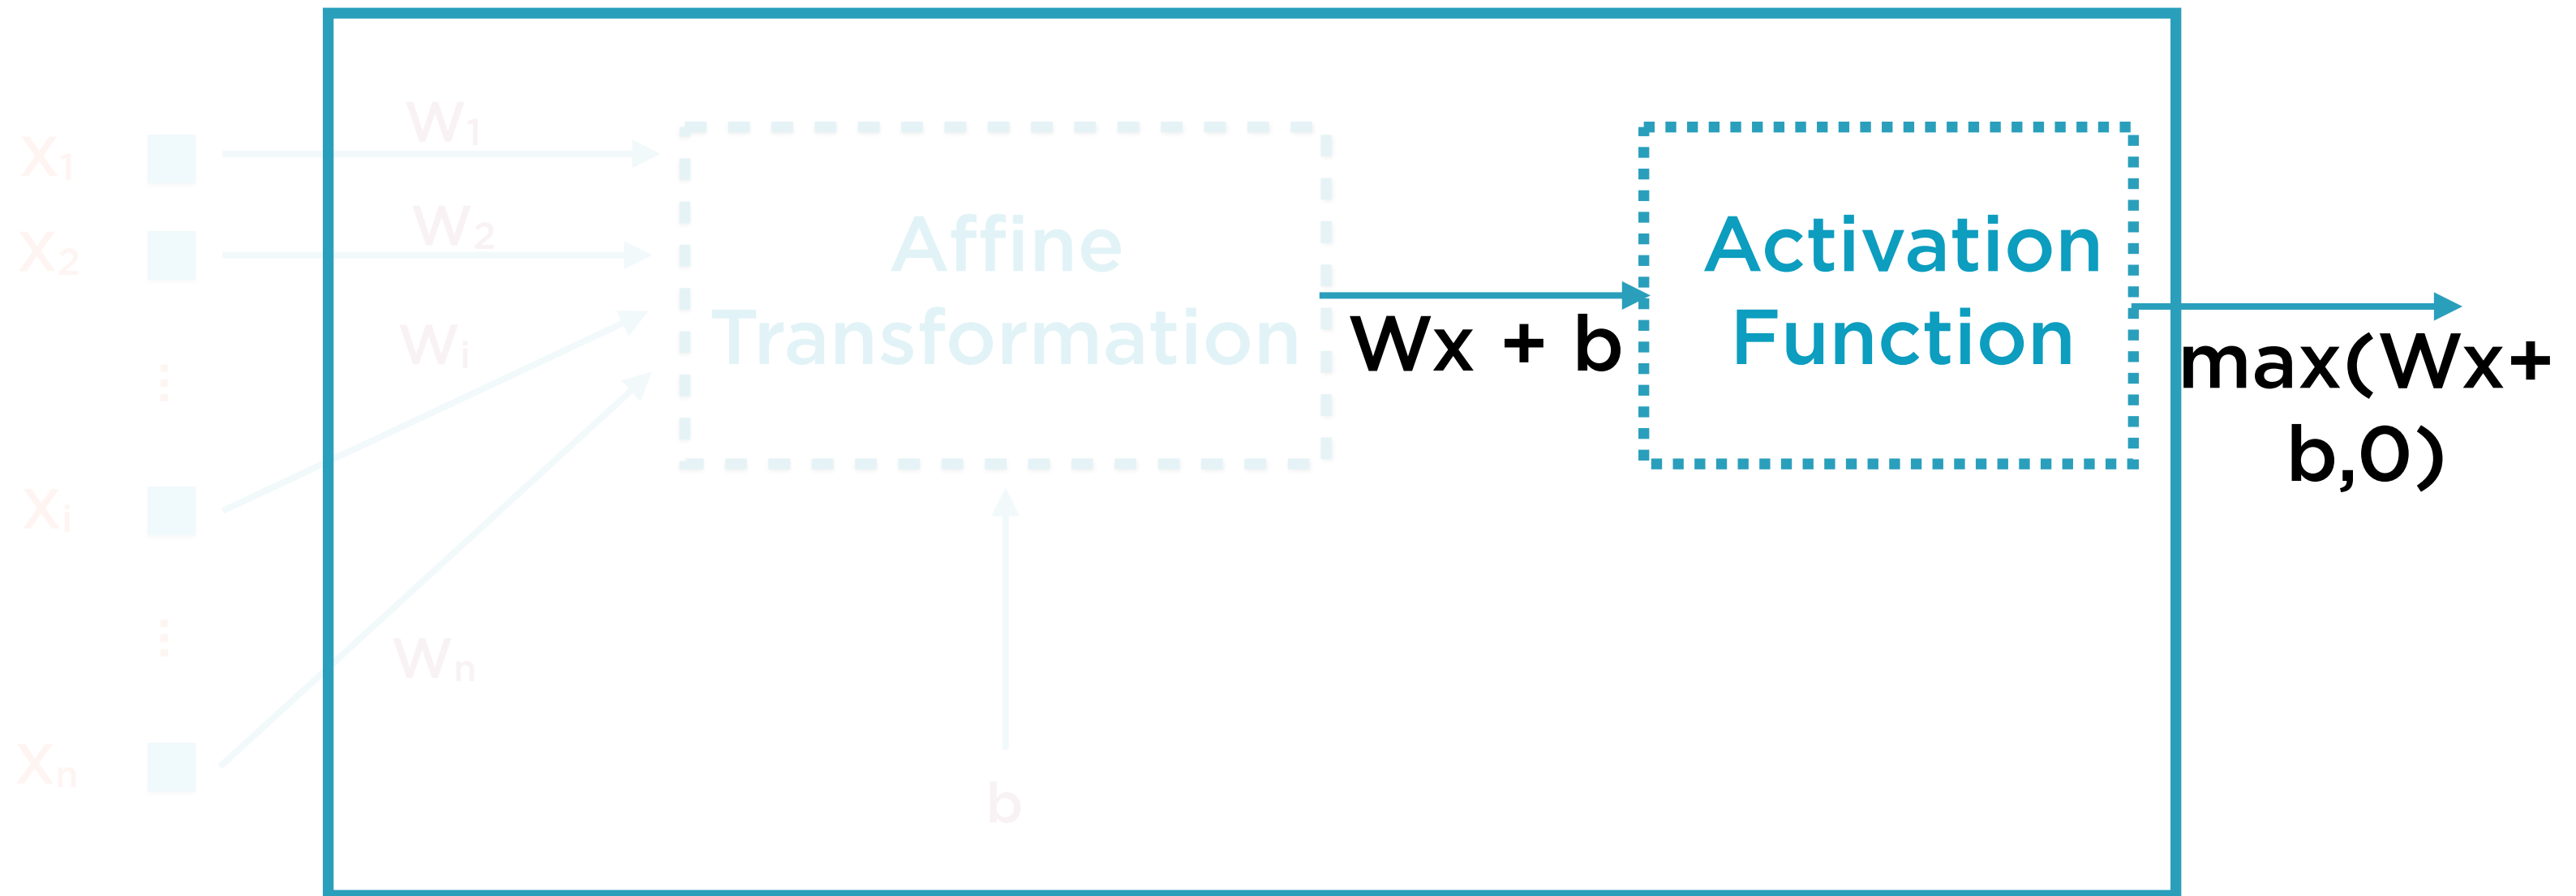
The affine transformation alone can **only** learn **linear** relationships between the inputs and the output

# Affine Transformation



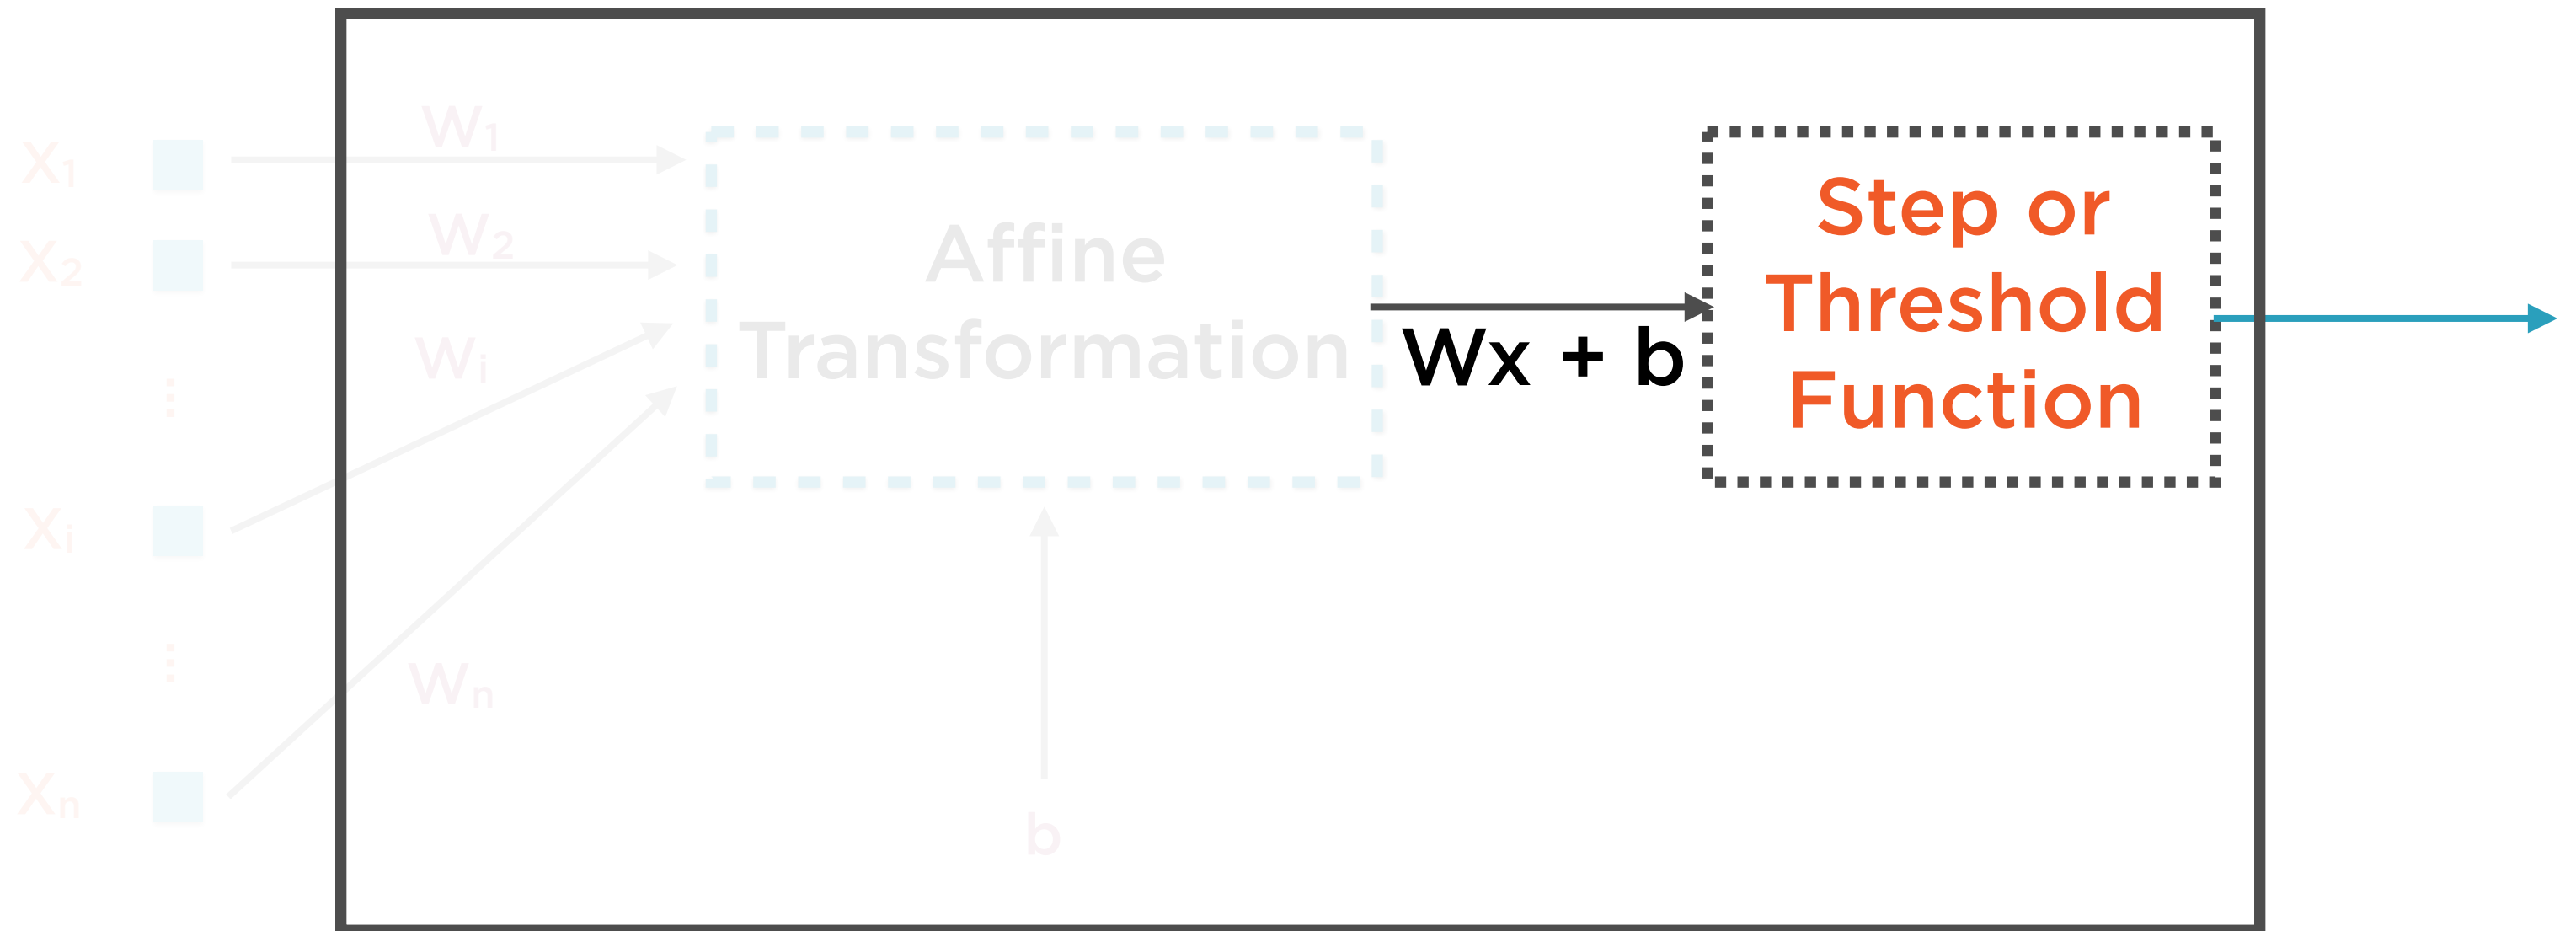
The affine transformation is just a weighted sum with a bias added:  $w_1x_1 + w_2x_2 + \dots + w_nx_n + b$

# Activation Function




**A function which helps discover non-linear relationships**

# Activation in a Perceptron



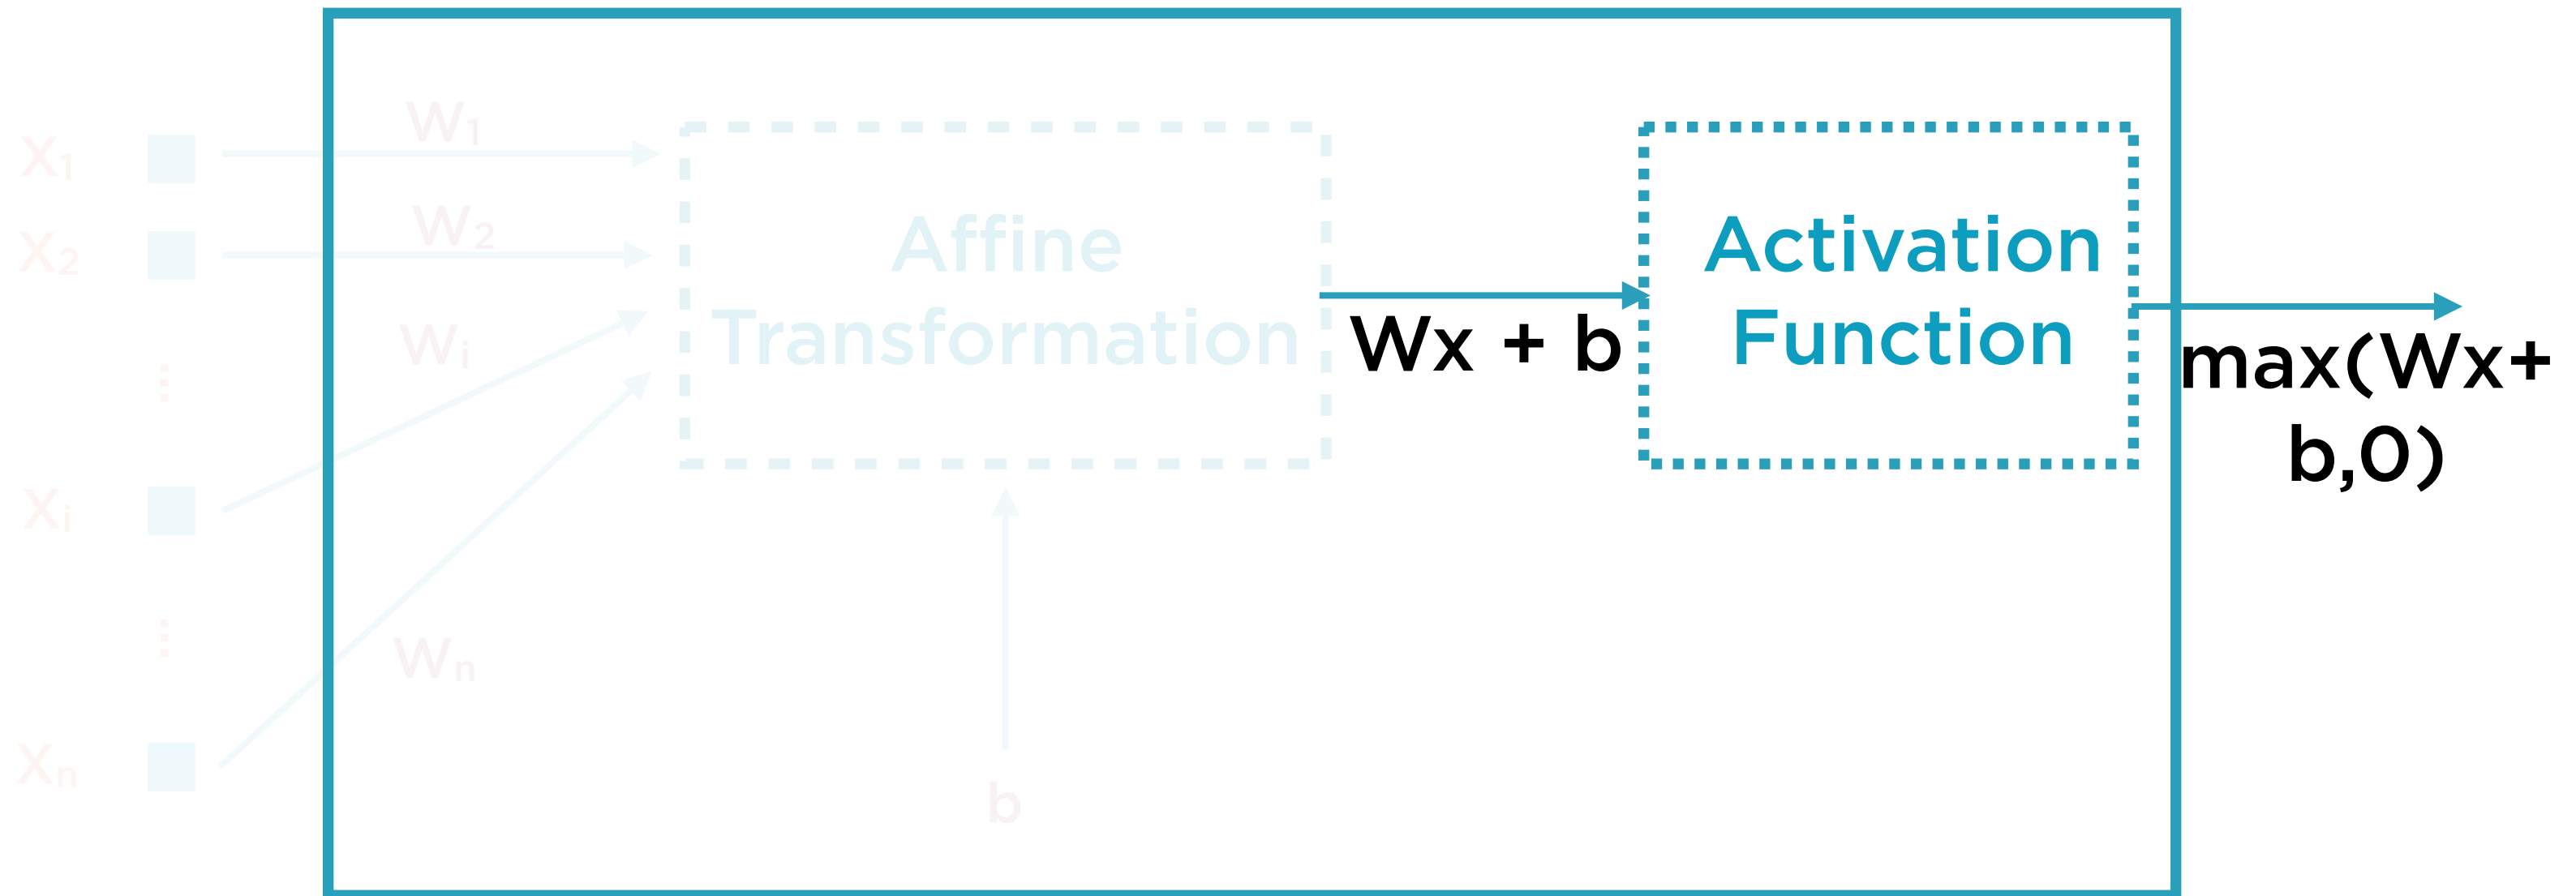
**This combination allowed the perceptron to only work with linearly separable data**

# Linearly Separable Data



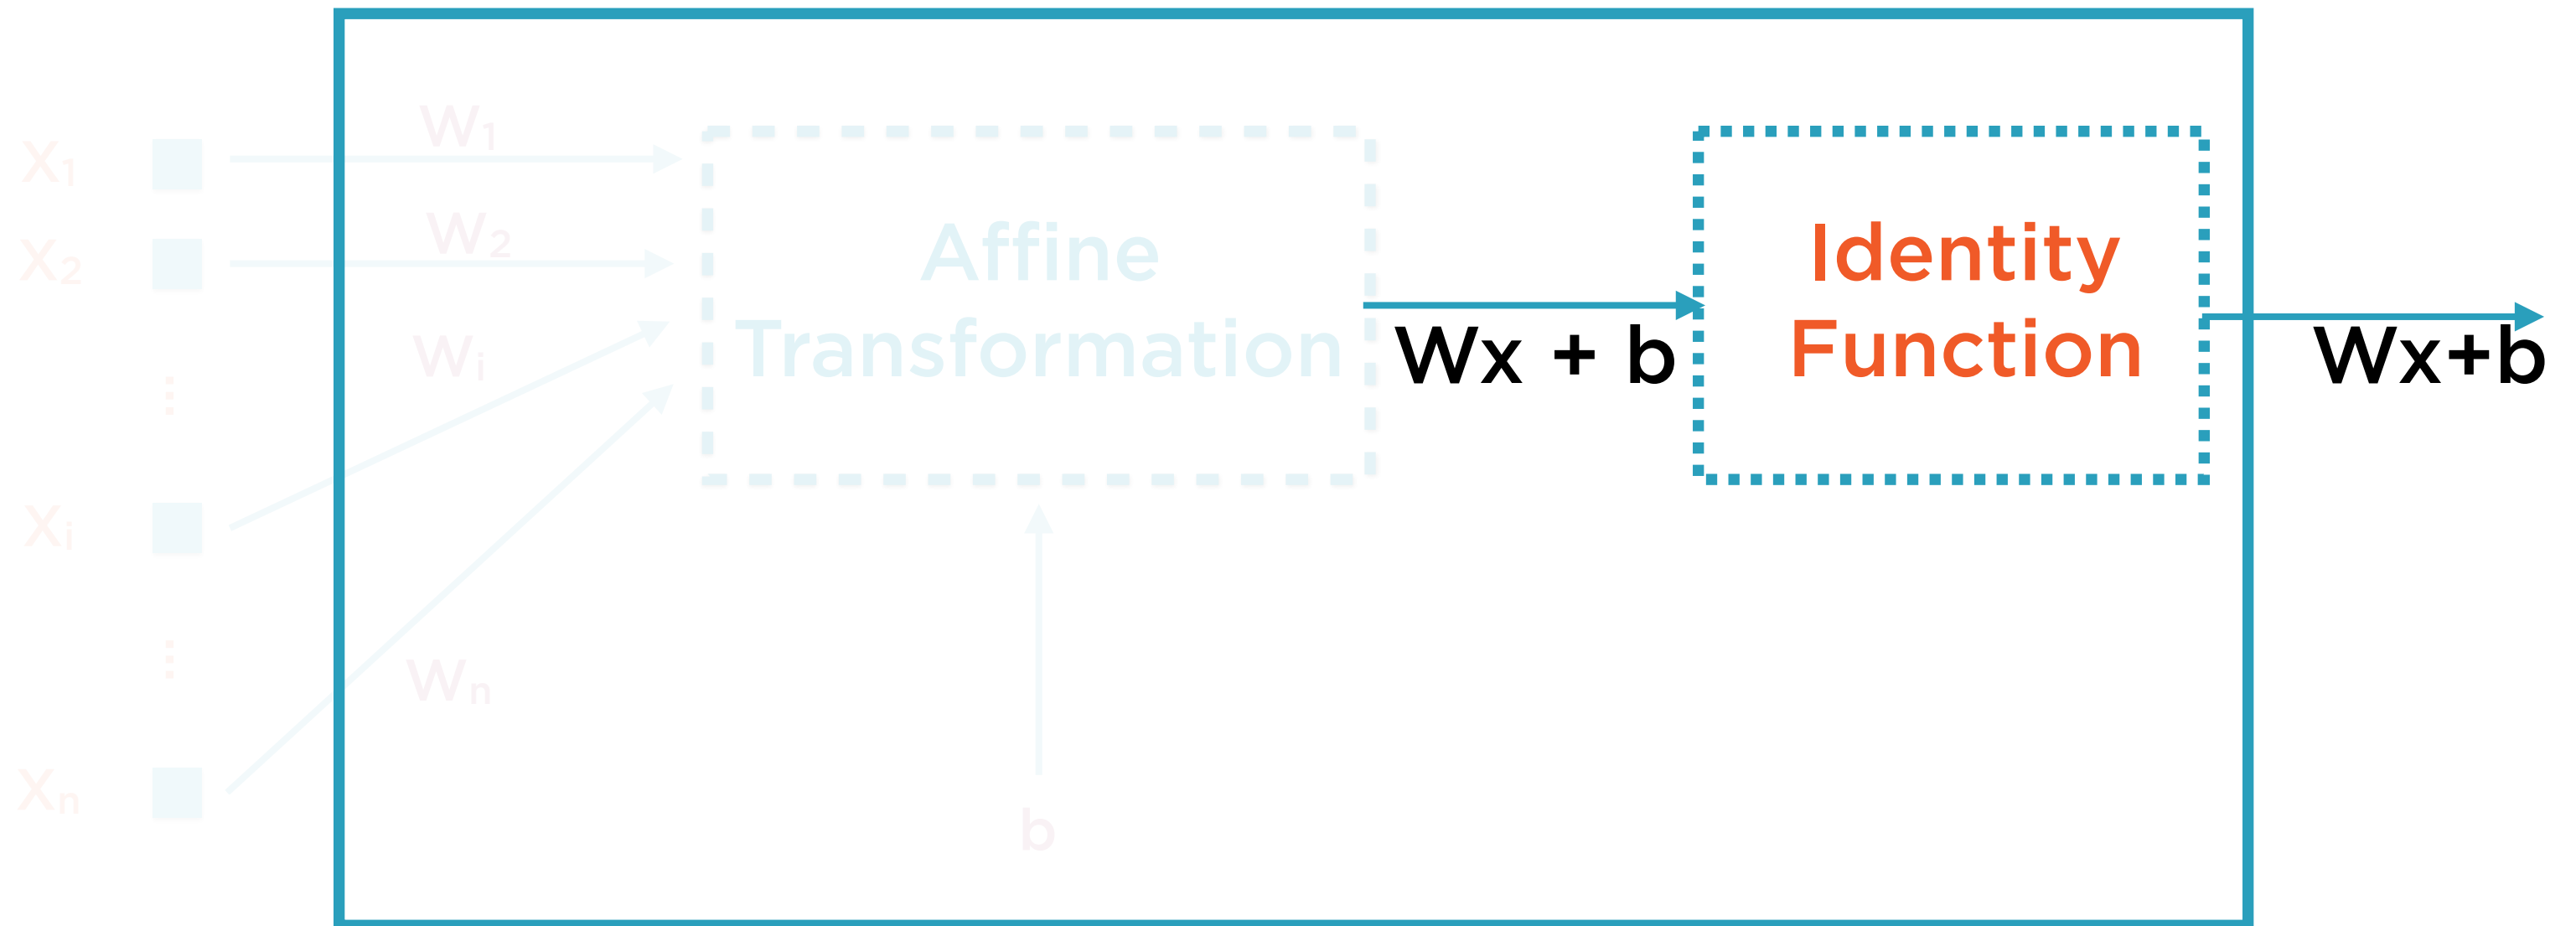
**Linear boundary between classes**

# Activation Function



**A neuron works with many more activation functions which help it learn more complex relationships**


# Linear Neuron



**When the activation function is the identity function, the neuron is often referred to as a linear neuron**




# Activation Function




The **combination** of the affine transformation and the activation function can **learn any arbitrary relationship**

# Common Activation Functions


**ReLU**




**logit**



**tanh**




**step**




# Active Region


ReLU




logit



tanh




step




Notice how activation functions have a gradient, this gradient allows them to be sensitive to input changes

# Saturation


ReLU




logit



tanh




step



In order to train and adjust the weights of the neural network the activation functions should operate in their active region

# Neuron as a Learning Unit



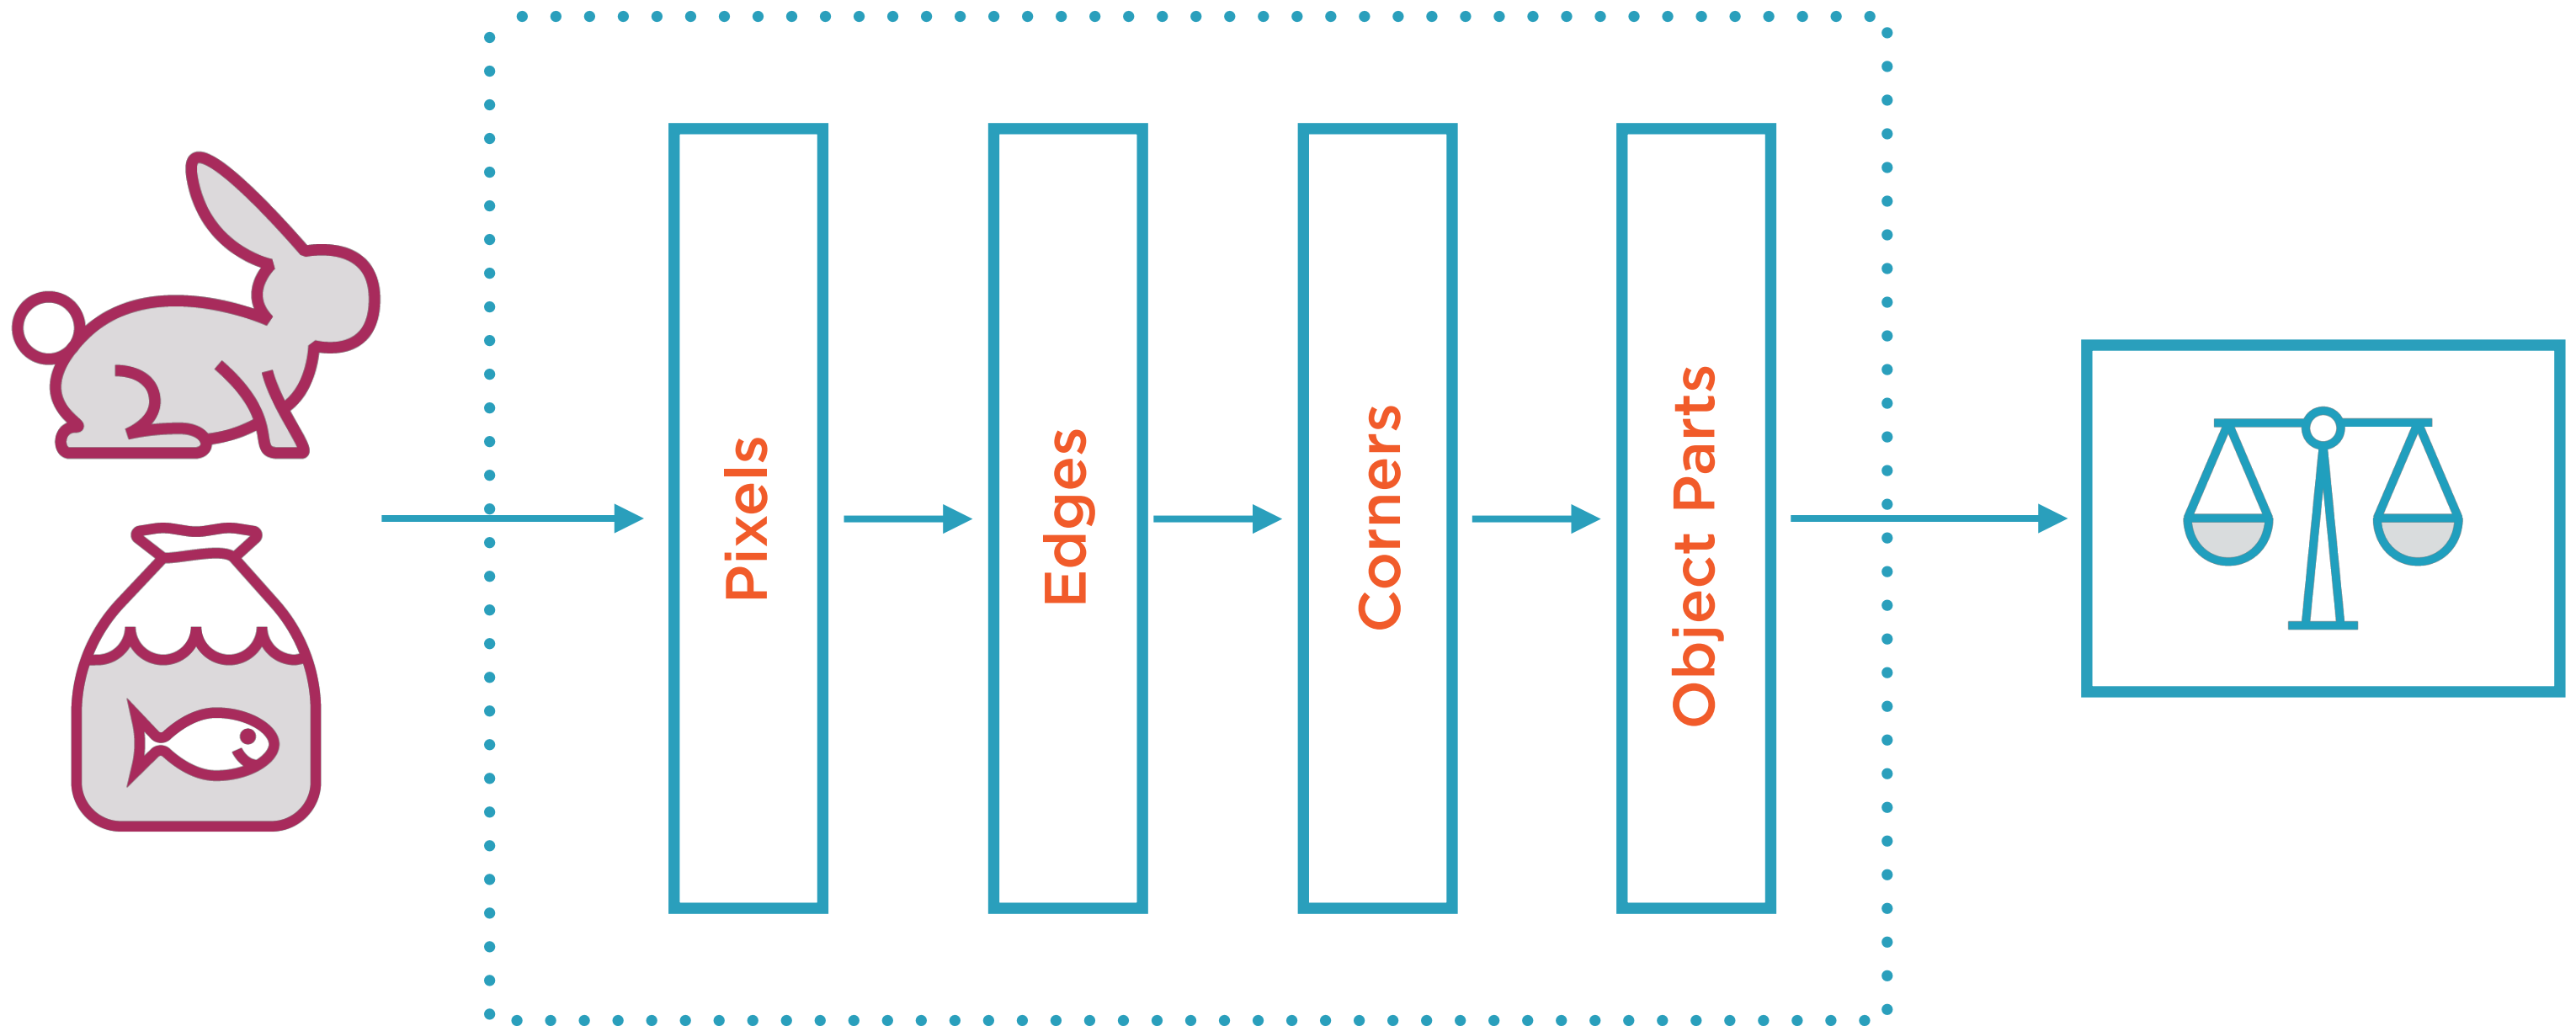
Many of these simple neurons arranged in layers can do magical stuff

# Multi-layer Perceptrons and Neural Networks

---

# Multi-layer Perceptron ~ Feed-forward Neural Network

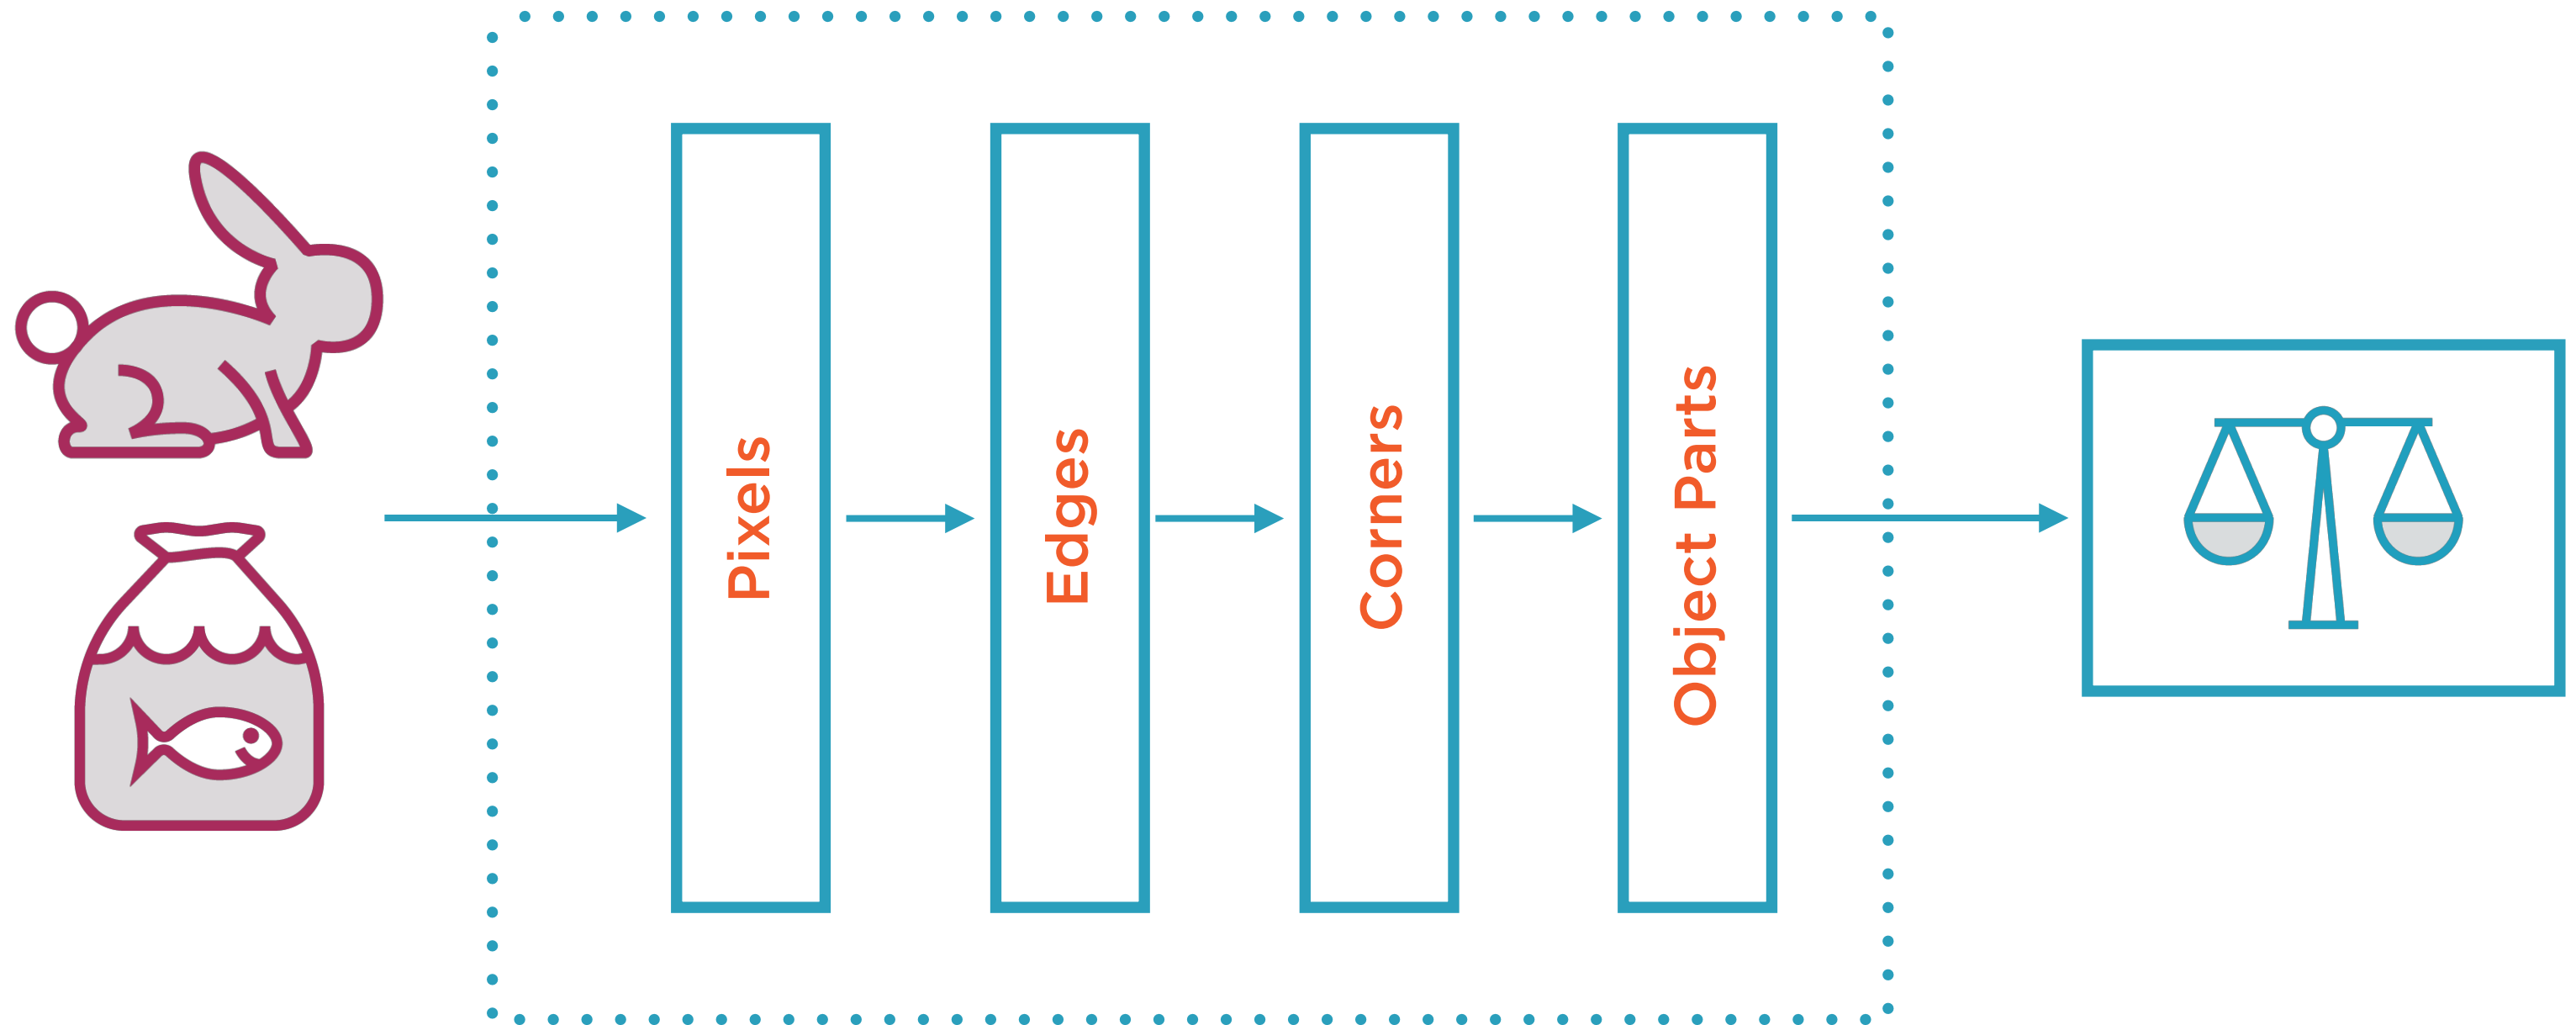
# Neural Network



**Comprised of neurons i.e. active learning units arranged in layers**




# Layers in a Neural Network




Groups of neurons that perform similar functions are aggregated into layers

# Network of Neurons



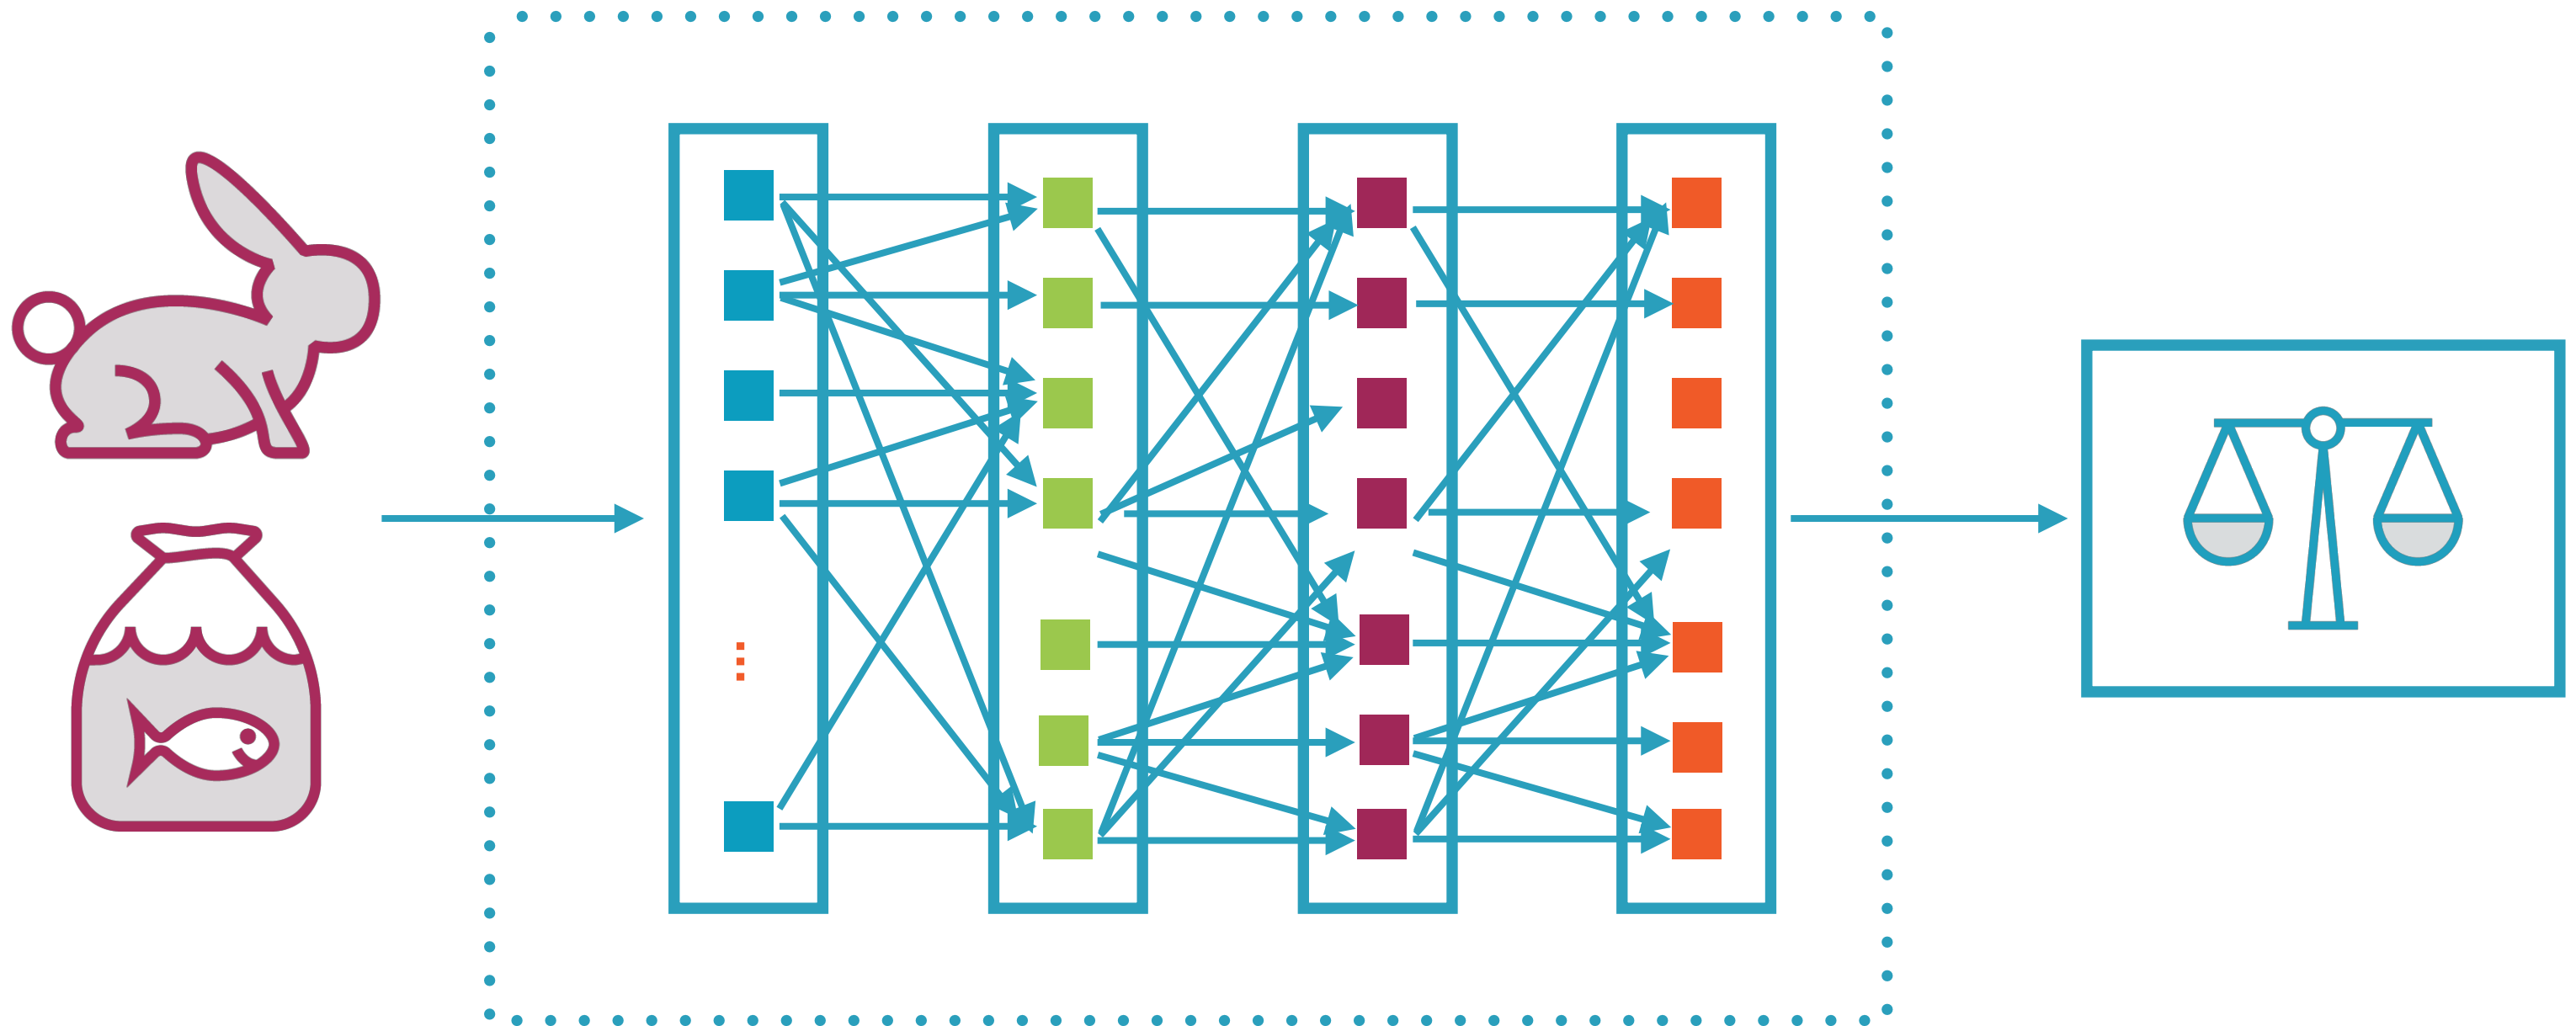
**Each layer consists of individual interconnected neurons**

# Network of Neurons




**Neurons receive input from  
neurons in the previous layer**

# Network of Neurons




**And pass their output on to  
neurons in the next layer**

# Feed-forward Networks




# Feed-forward Networks




**Feed-forward networks: Information moves forward through the layers**

# Feed-forward Networks




**No connection exists between  
neurons in the same layer**


# Deep Learning with Neural Networks



**Directed computation graphs “learn” relationships between data**



**The more complex the graph, the more relationships it can “learn”**



**“Deep” Learning: Depth of the computation graph**



$$y = Wx + b$$

---

## “Learning” Regression

**Regression can be reverse-engineered by a single neuron**

```
def doSomethingReallyComplicated(x1, x2...):  
    ...  
    ...  
    ...  
    return complicatedResult
```

---

## “Learning” Arbitrarily Complex Functions

**Adding layers to a neural network can “learn” (reverse-engineer) pretty much anything**

# Training a Neural Network

---

# Neurons




**The nodes in the computation graph are simple entities called neurons**

**Each neuron performs very simple operations on data**

**The neurons are connected in very complex, sophisticated ways**

# Neural Network




**The complex interconnections between simple neurons**

**Different network configurations => different types of neural networks**

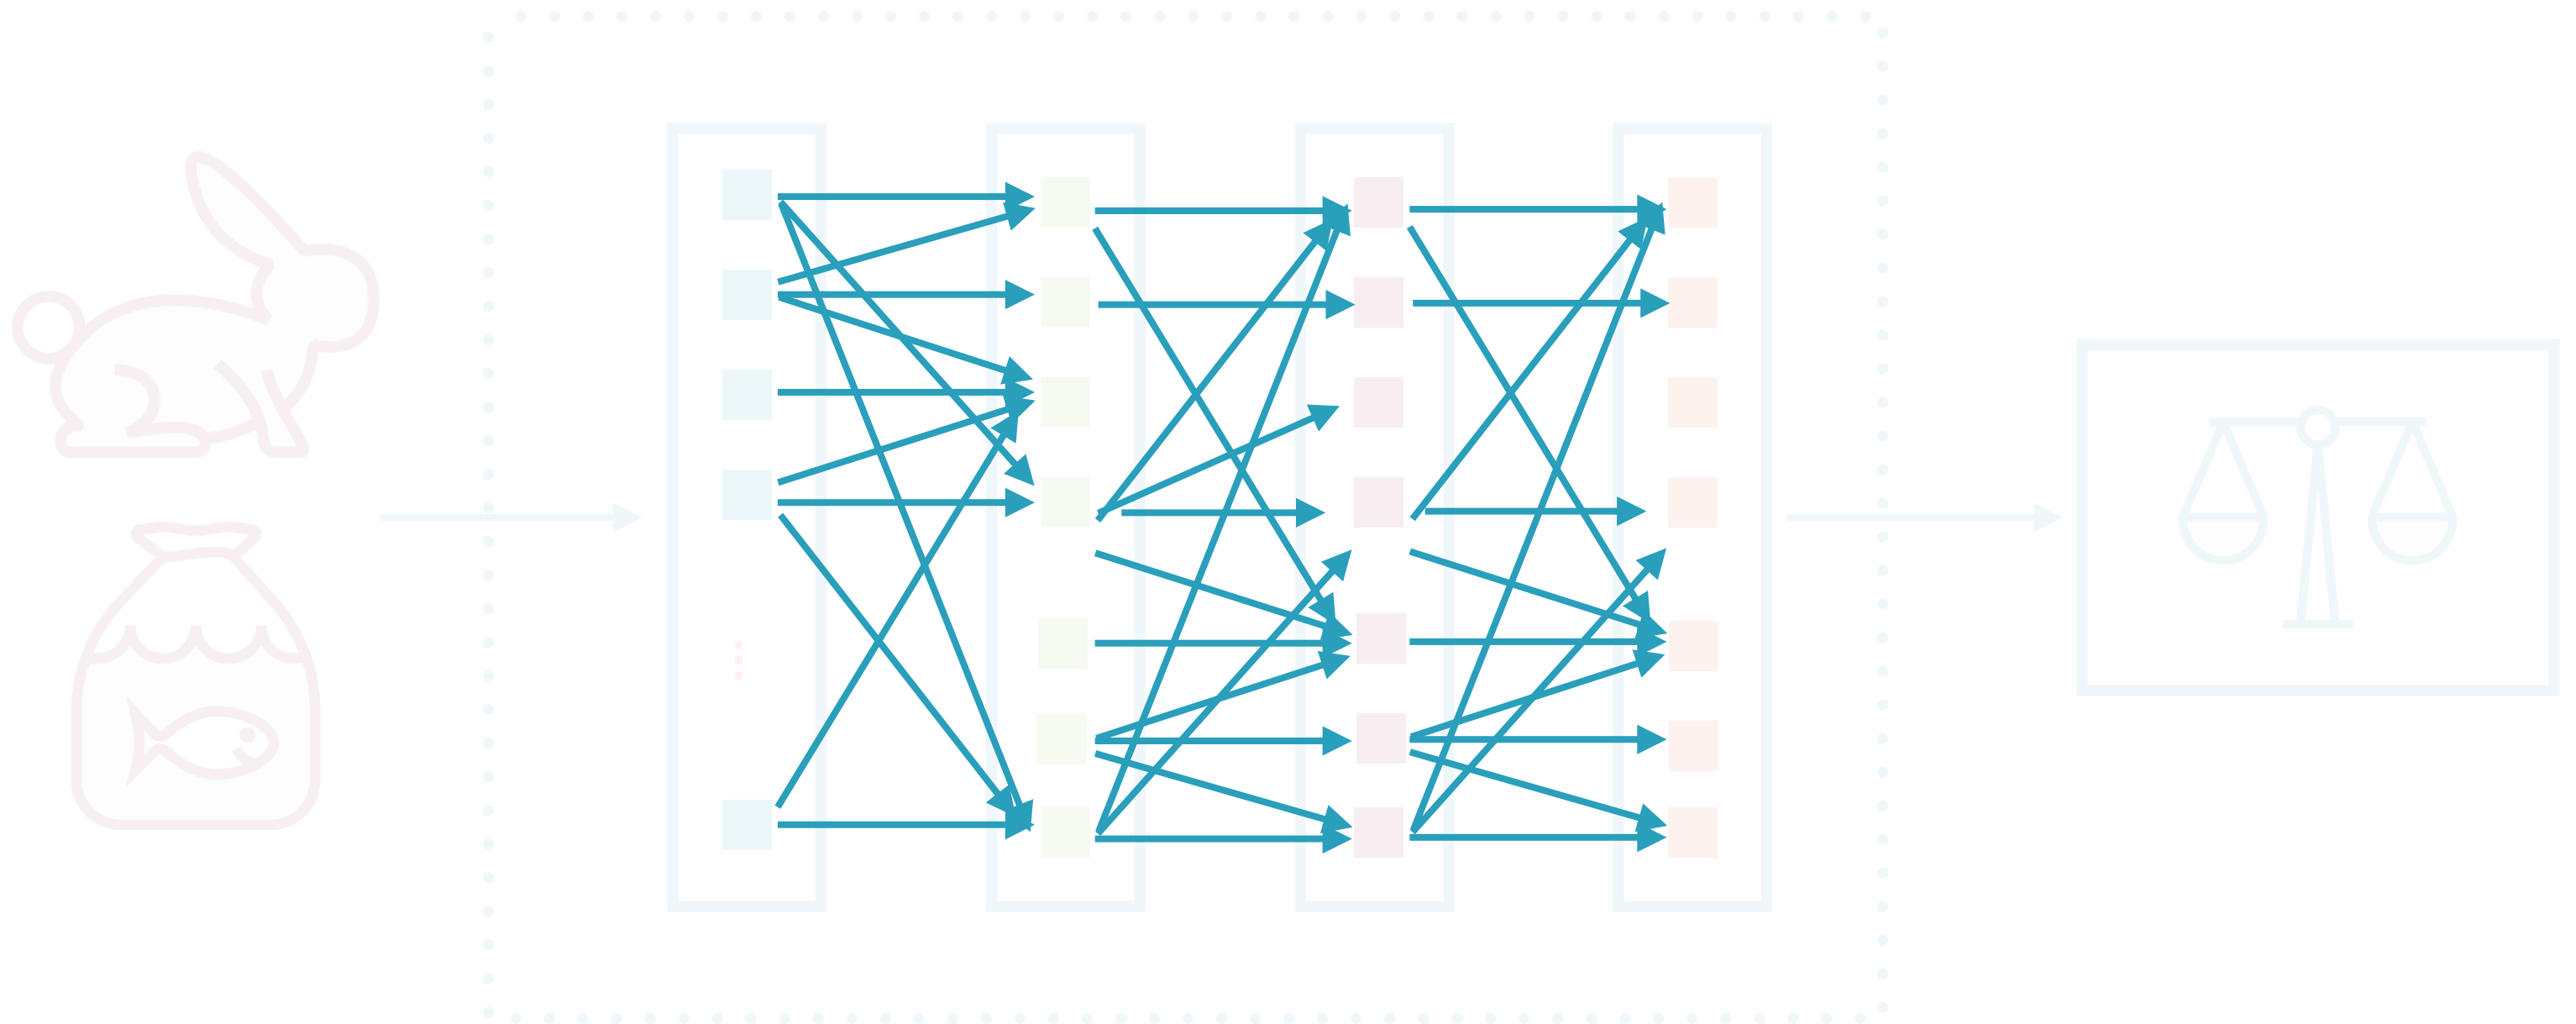
- Convolutional
- Recurrent

# Neural Network




**Groups of neurons that perform similar functions are aggregated into layers**

# Complex Interconnections



**Neurons in a neural network can be connected in very complex ways...**

# Neural Network

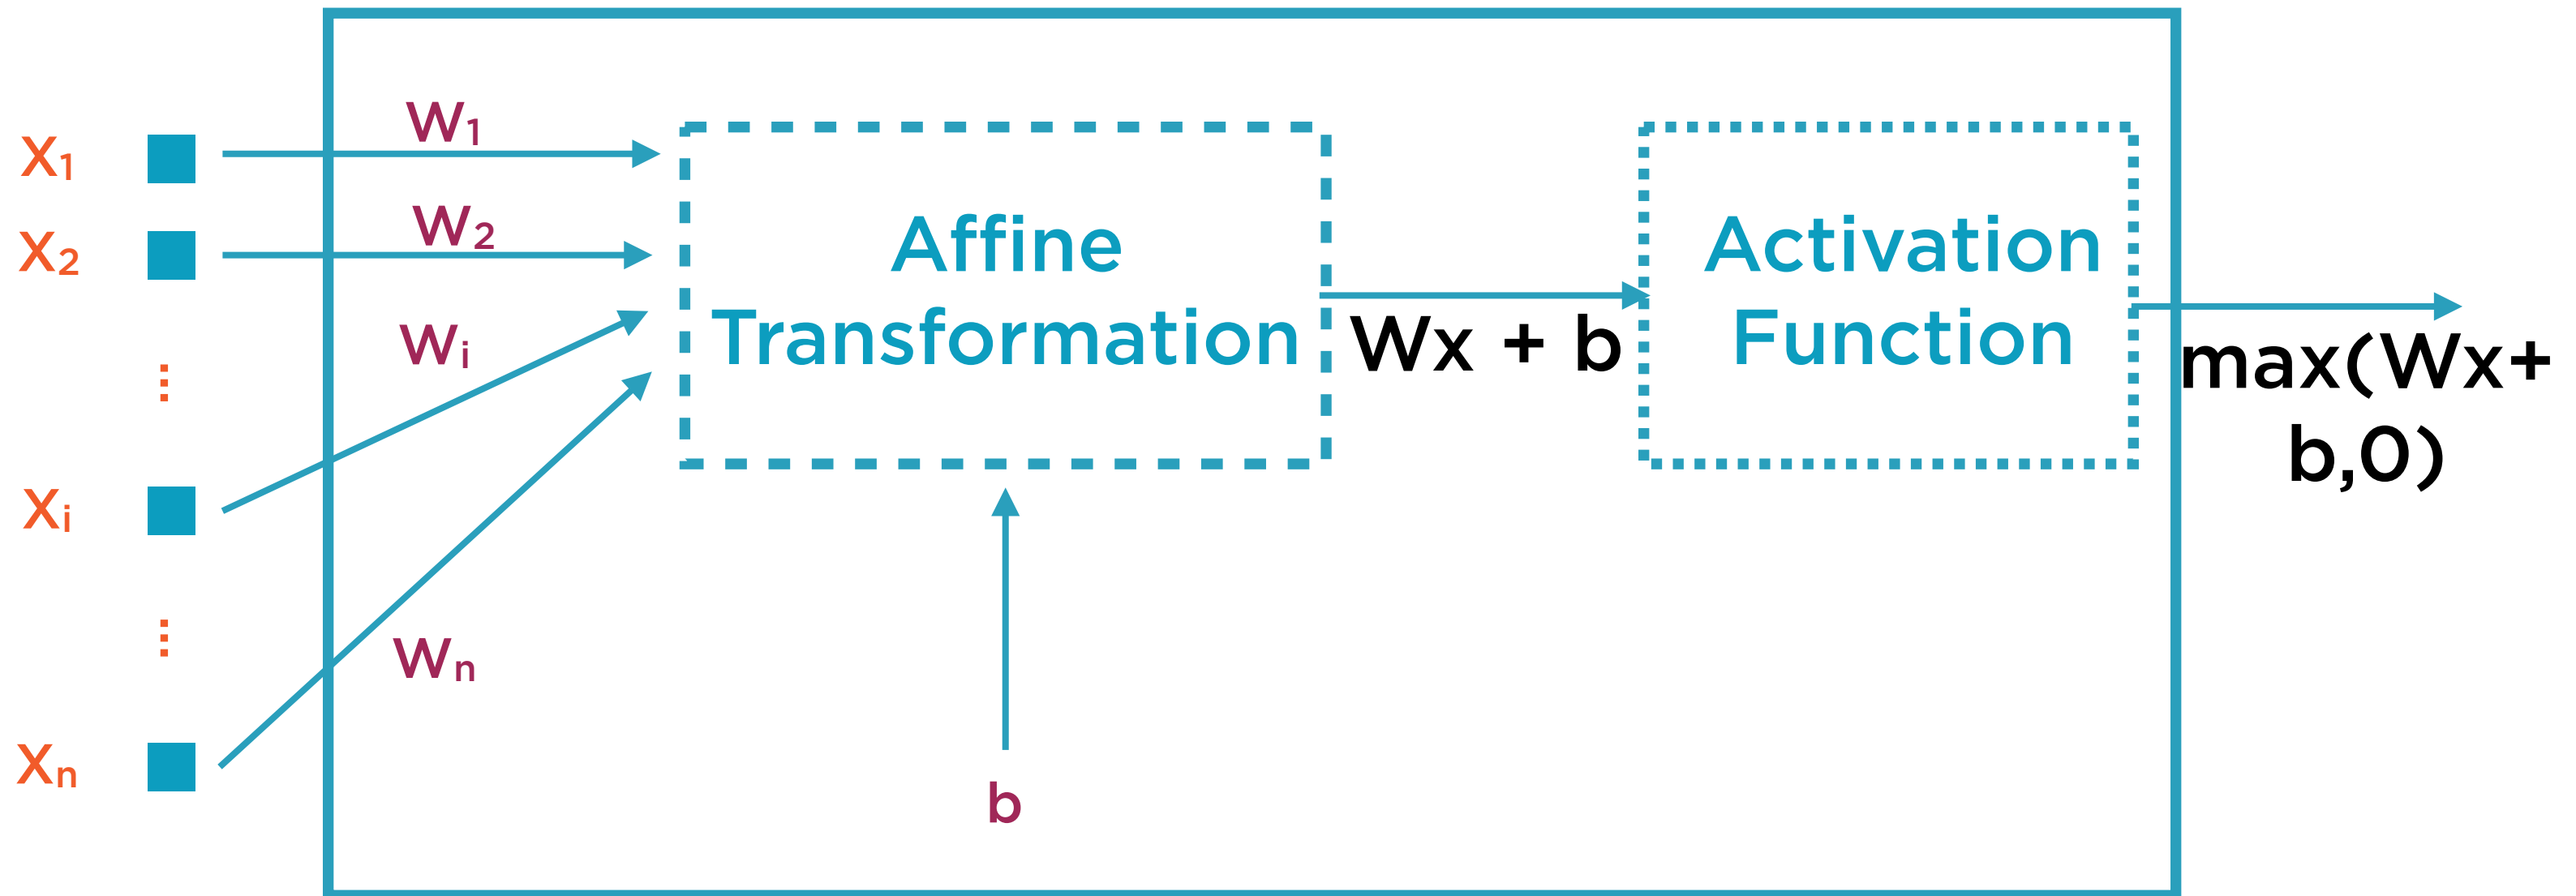


Neurons in a neural network can be connected in very complex ways...

...But each neuron only applies **two** simple functions to its inputs

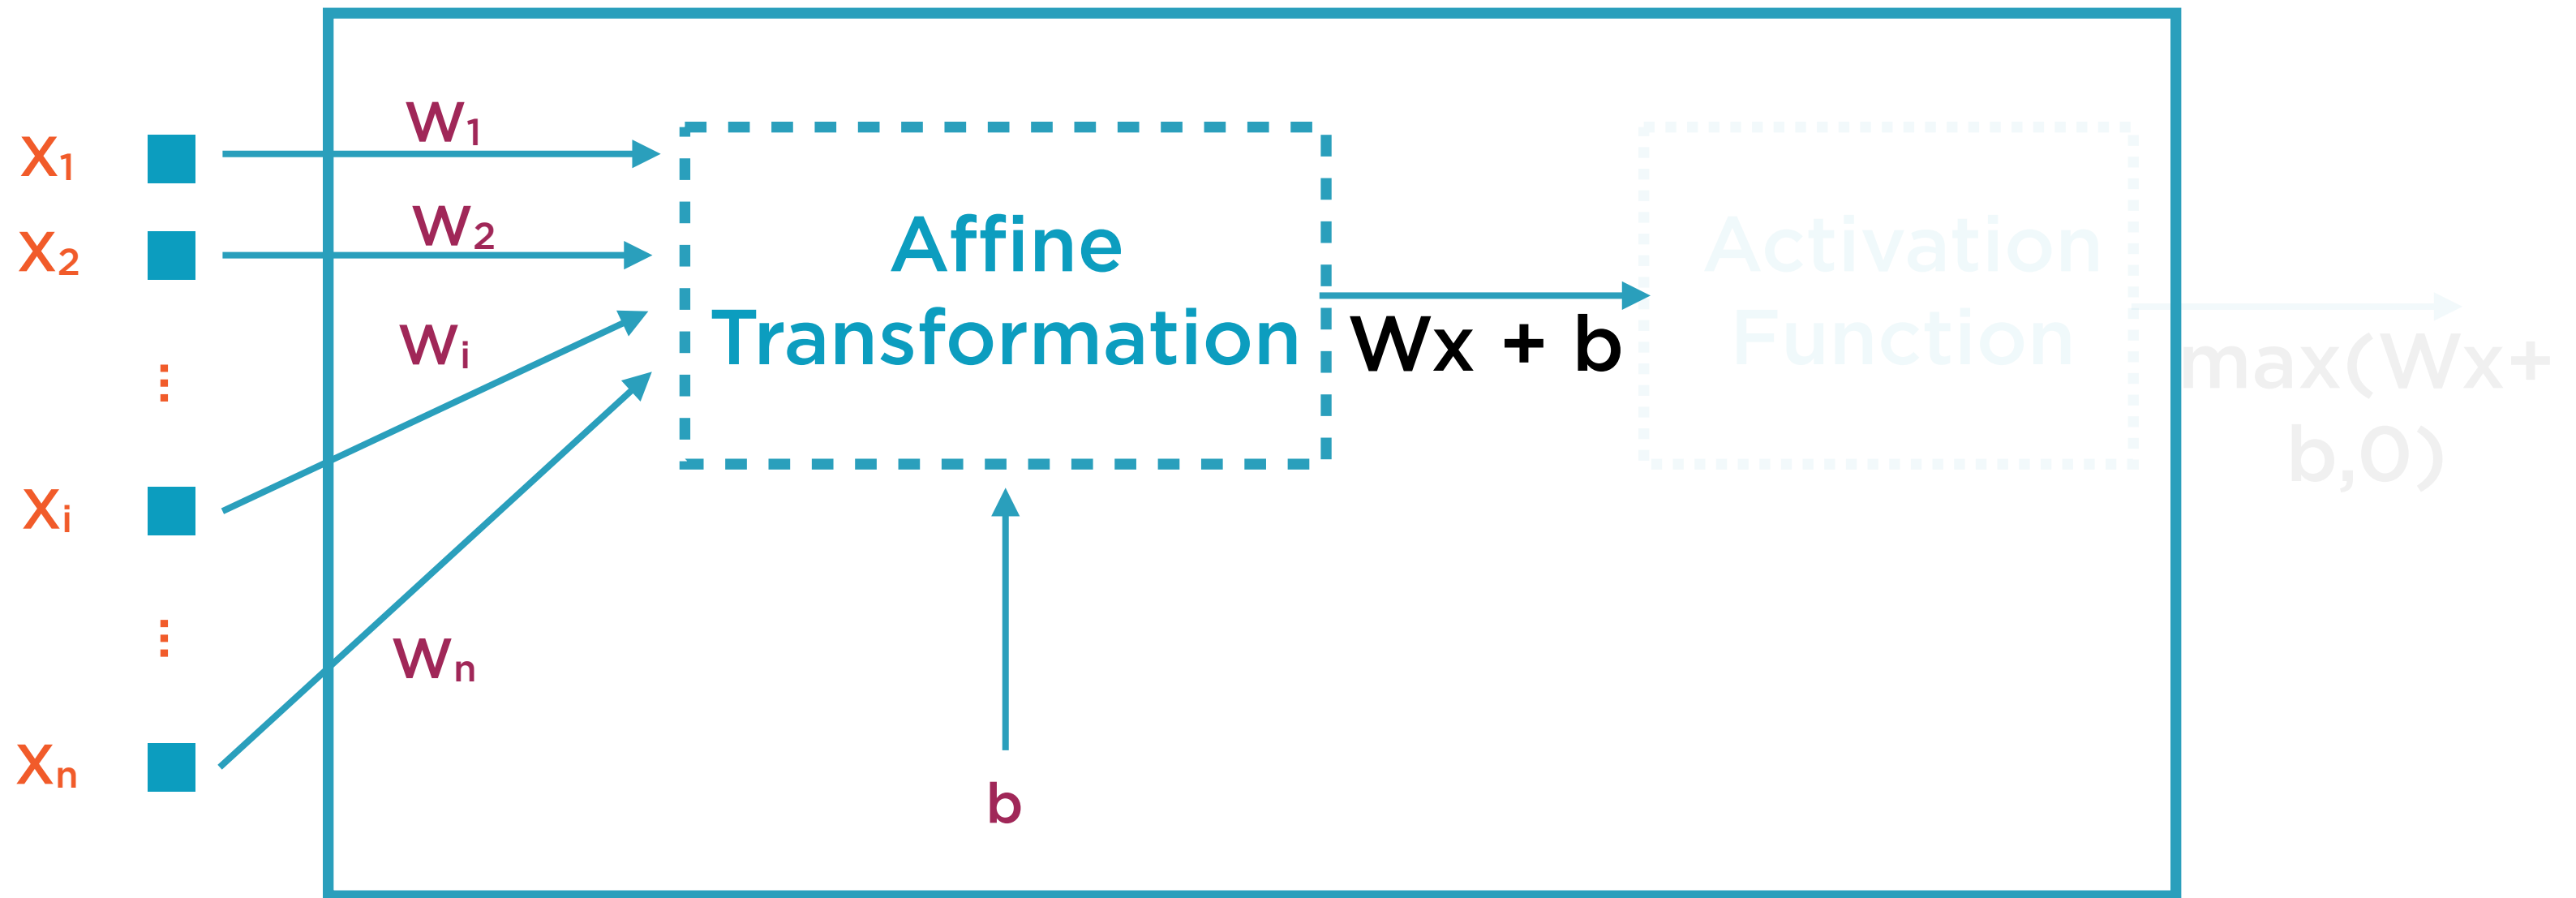


# Operation of a Single Neuron




**Each neuron only applies two simple functions to its inputs**

# Operation of a Single Neuron



**Where do the values of  $W$  and  $b$  come from?**




Finding the “best” values of  $W$  and  $b$  for each neuron is crucial

The “best” values are found using the cost function, optimizer and corpus


And the process of finding them is called the **training process**

Training a neural network happens using **gradient descent**


# Gradient Descent




# Gradient Descent



# Gradient Descent




# Gradient Descent




# Gradient Descent

Converging on the “best”  
value using an  
optimization algorithm





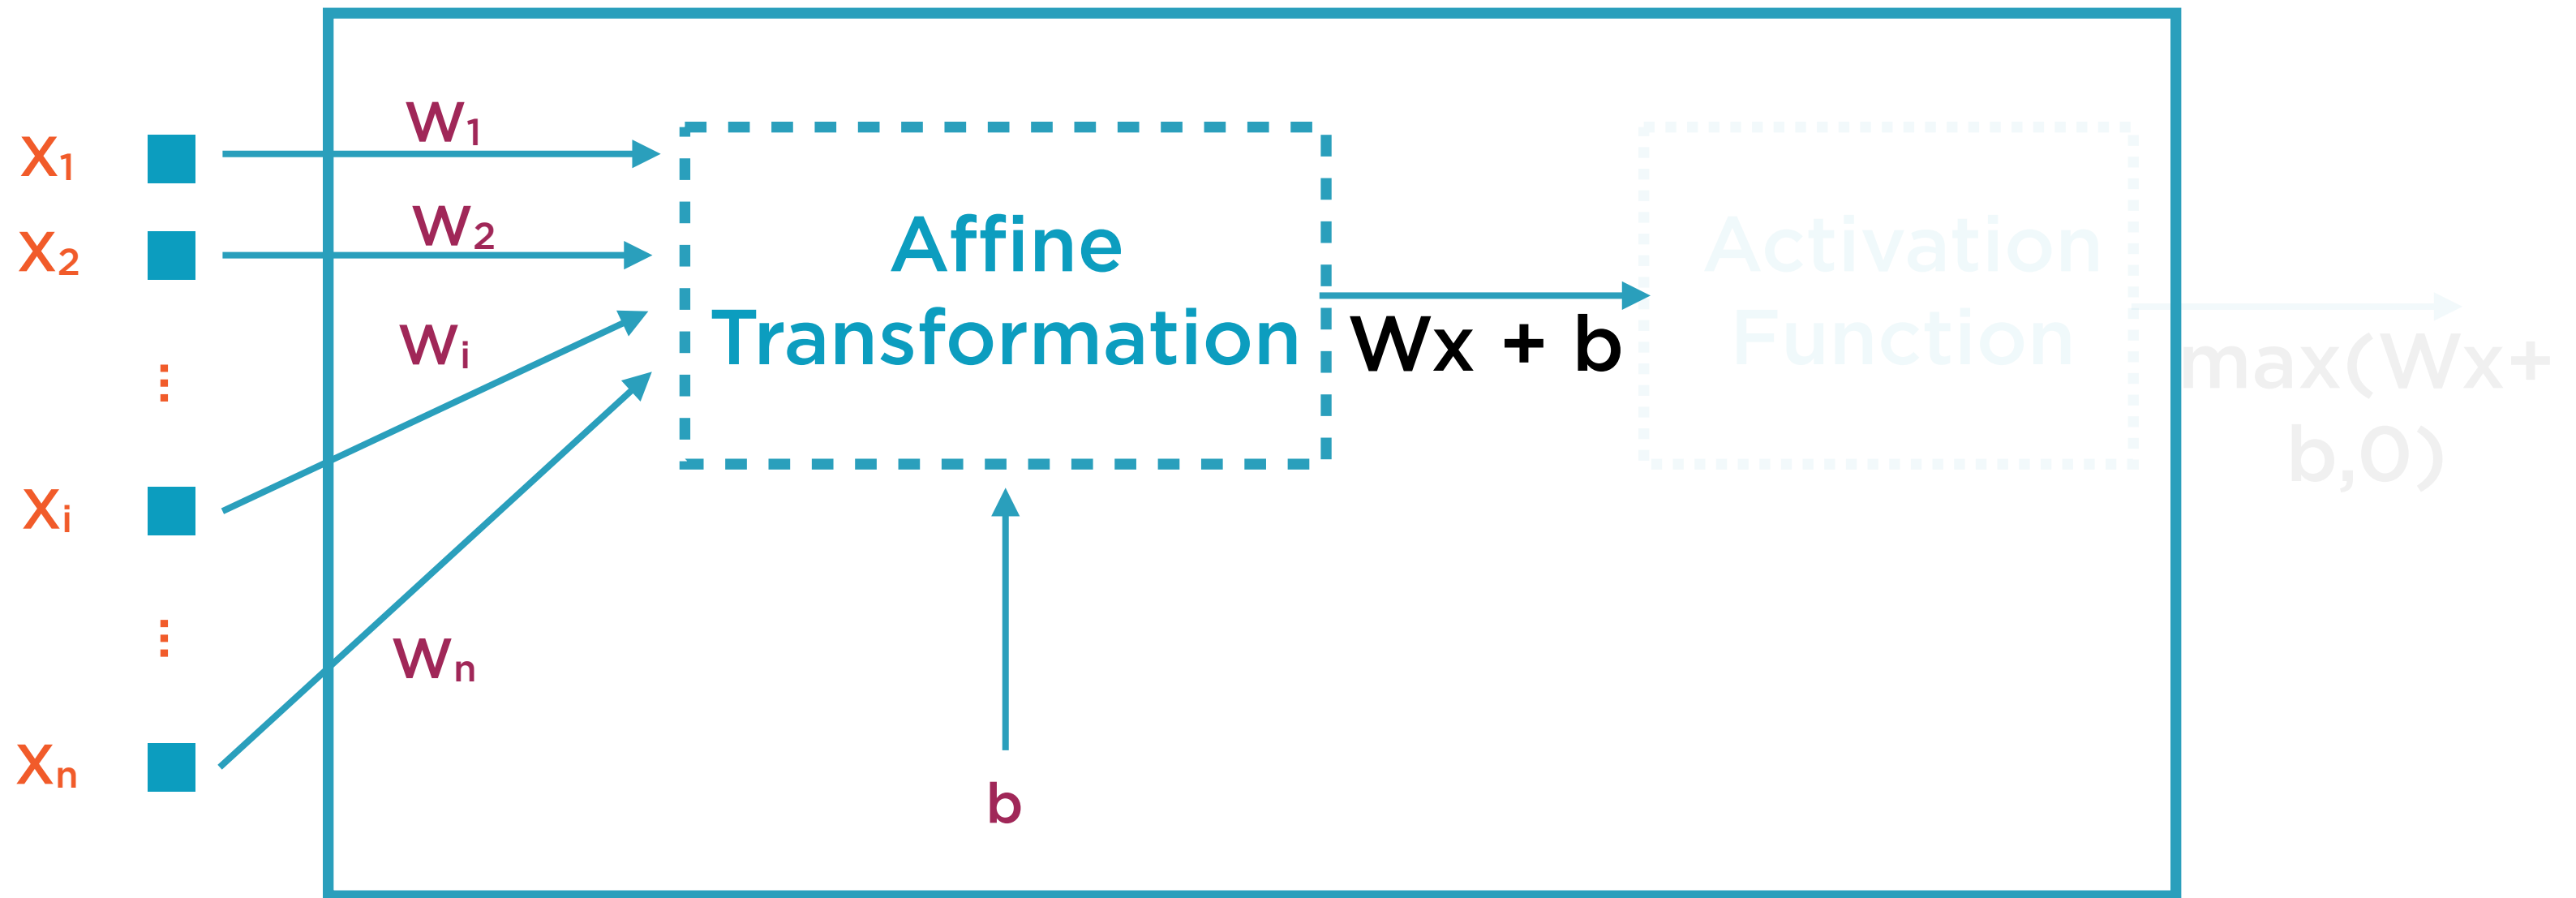


During training, the output of deeper layers may be “fed back” to find the best  $W$ ,  $b$

This is called **back propagation**


Back propagation is the standard algorithm for training neural networks

# Operation of a Single Neuron




**The training algorithm will use the weights to tell the neuron which inputs matter, and which do not...**

# Operation of a Single Neuron



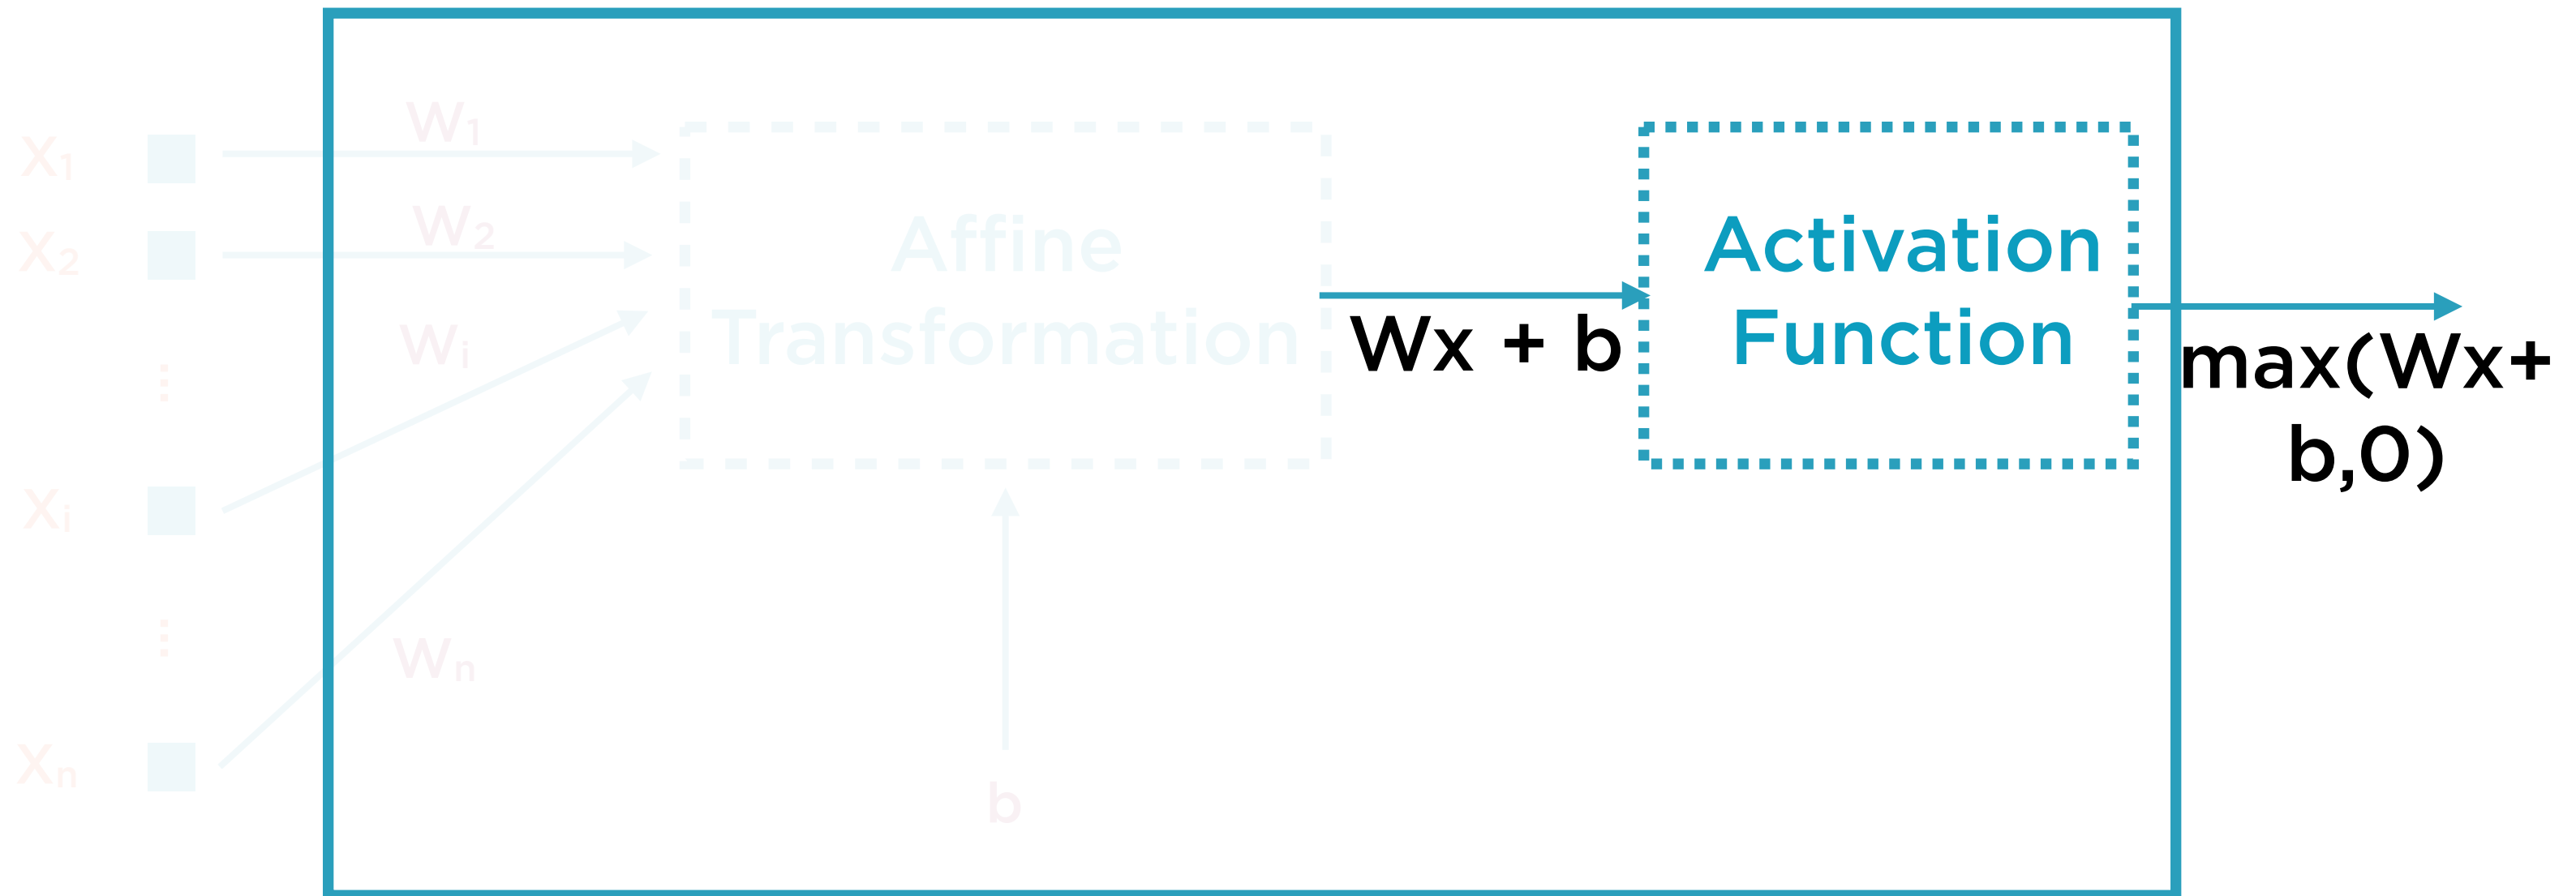
...and apply a corrective bias if needed

# Operation of a Single Neuron



**The linear output can only be used to learn linear functions, but we can easily generalize this...**

# Operation of a Single Neuron



**The Activation function is a non-linear function, very often simply the  $\max(0, \dots)$  function**



**The output of the affine transformation is chained into an activation function**



The activation function is needed for the neural network to predict non-linear functions



The most common form of the activation function is the ReLU

ReLU : Rectified Linear Unit


$$\text{ReLU}(x) = \max(0, x)$$



# Overfitting and Underfitting


---

# Connecting the Dots




Challenge: Fit the “best” curve through these points

# Good Fit?




A curve has a “good fit” if the distances of points from the curve are small

# Connecting the Dots




We could draw a pretty complex curve

# Connecting the Dots




We can even make it pass through every single point

# Connecting the Dots




But given a new set of points, this curve might perform quite poorly

# Connecting the Dots



The original points were “training data”, the new points are “test data”


# Overfitting



Great performance in training, poor performance in real usage




# Connecting the Dots



A simple straight line performs worse in training, but better with test data

# Overfitting



**Model has memorized the training data**

**Low training error**


**Does not work well in the real world**

**High test error**

# Preventing Overfitting



**Regularization - Penalize complex models**




**Cross-validation - Distinct training and validation phases**



**Dropout (NNs only) - Intentionally turn off some neurons during training**

# Underfitting



**Model unable to capture relationships in data**

**Performs poorly on the training data**

**Model too “simple” to be useful**

# Summary

**scikit-learn support for neural networks**

**scikit-learn vs. deep learning frameworks**

**Perceptrons and neurons**

**Multi-layer perceptrons (MLPs) and neural networks**

**Training a neural network**