

Implementing Regression and Classification Using Neural Networks in scikit-learn



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Regression models using neural networks in scikit-learn

Mean square error (MSE) loss function

Using the MLPRegressor estimator to build regression models


Classification models using neural networks in scikit-learn

Cross-entropy loss function

Using the MLPClassifier estimator to build classification models

Regression Using Neural Networks


Neural Network Model



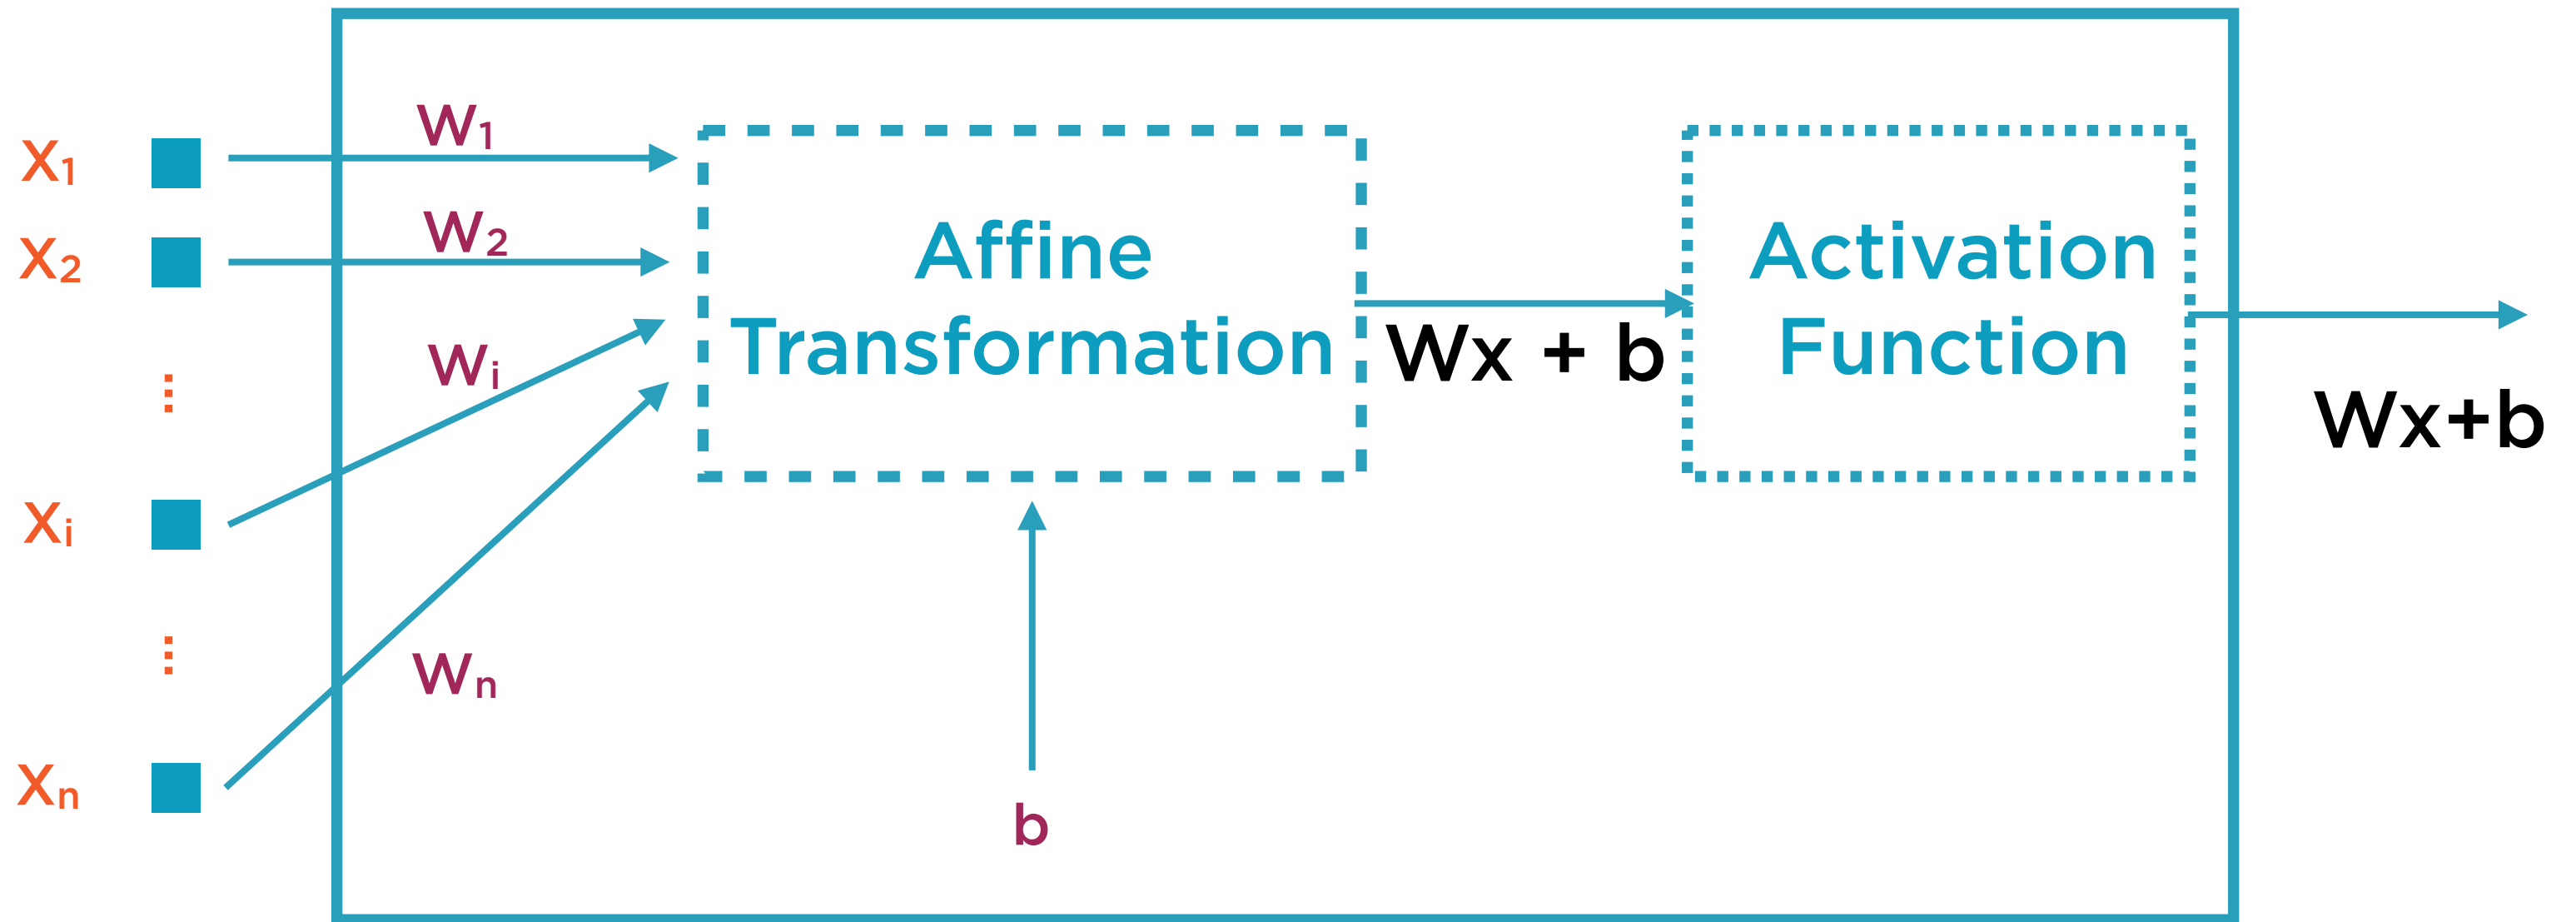
Network of interconnected layers

The **weights** and **biases** of individual neurons are determined during the **training** process

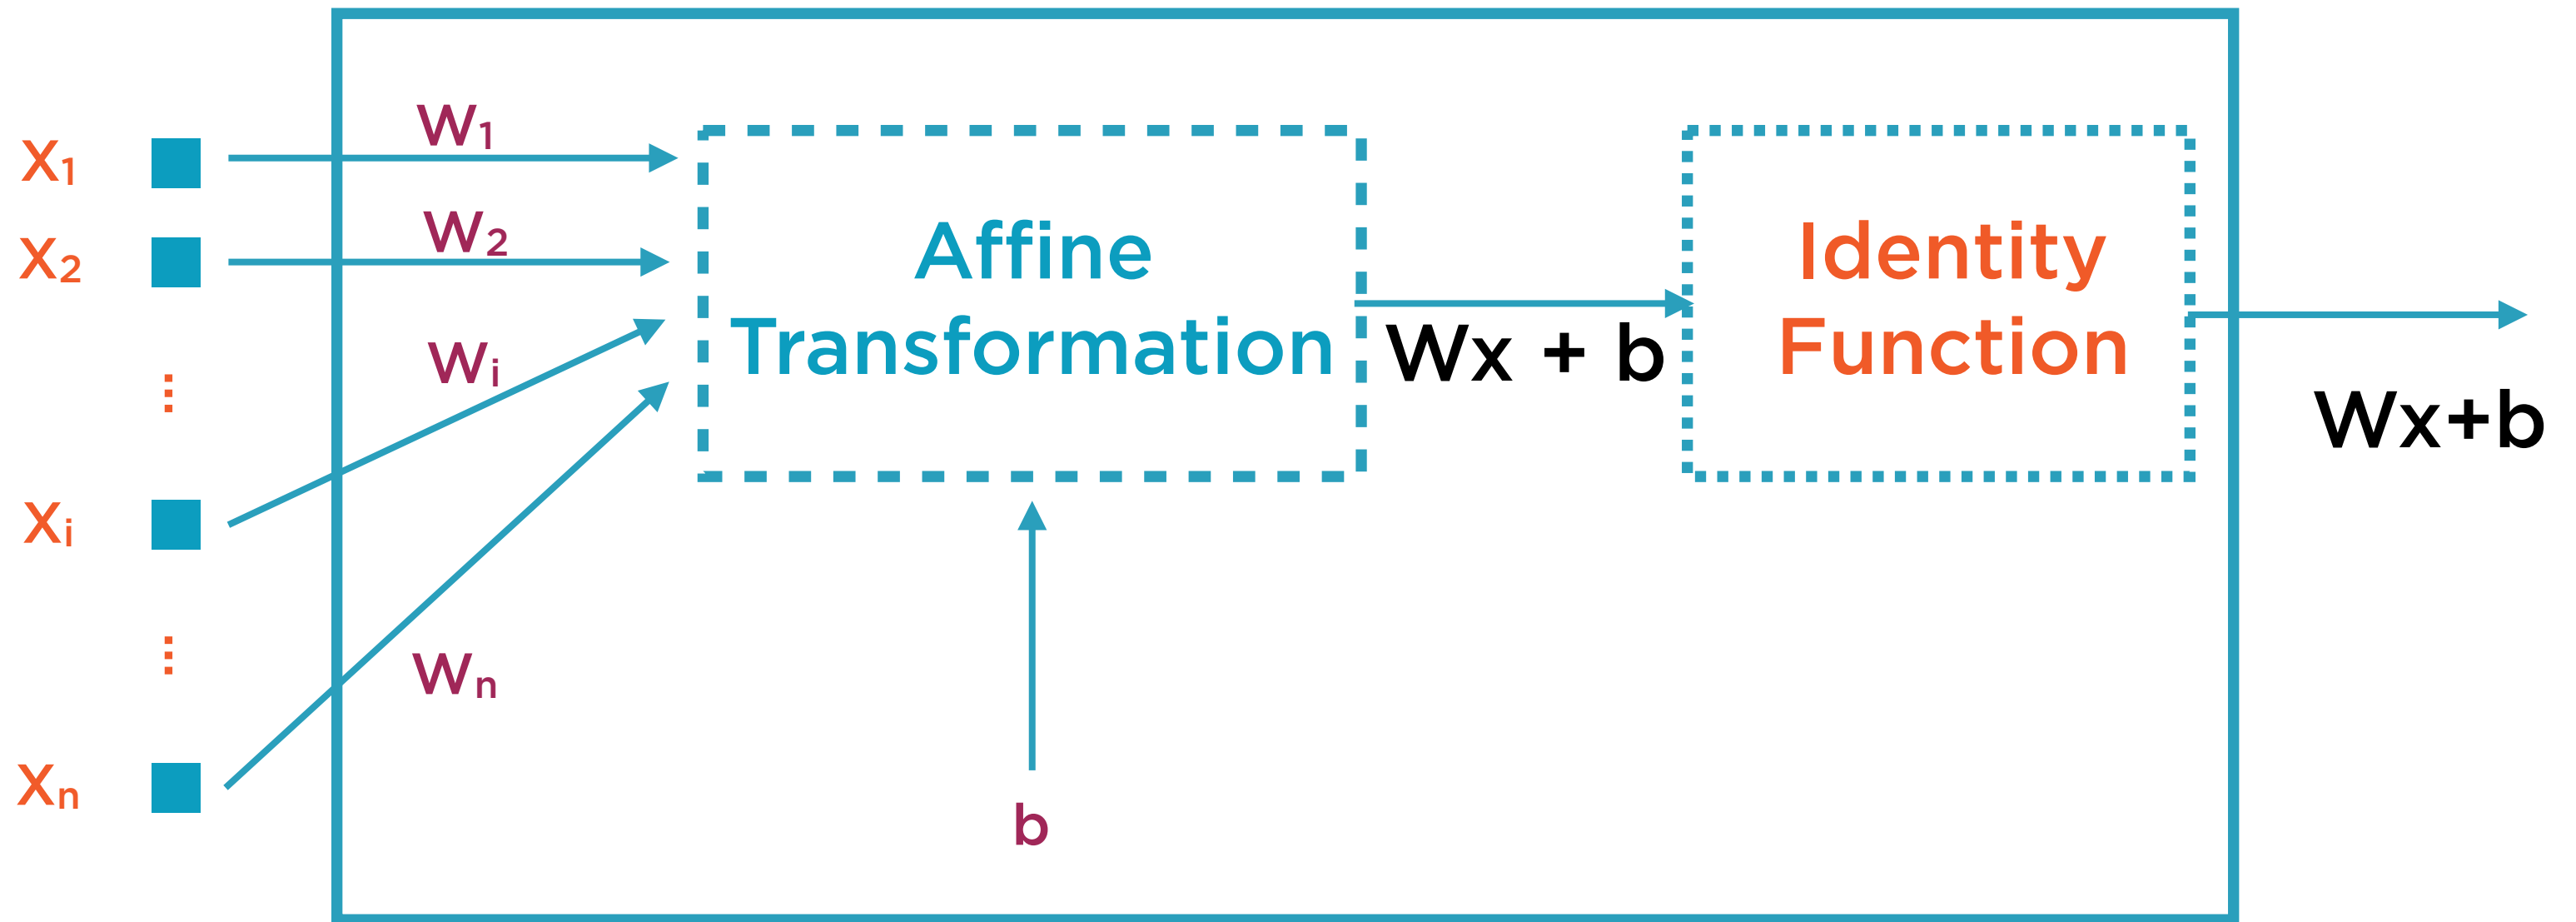
Regression: The Simplest Neural Network



Regression: The Simplest Neural Network




Regression: The Simplest Neural Network



Minimizing Mean Square Error




Y



Regression Line:
 $y = A + Bx$

X



The “best fit” line is called the
regression line

The actual training of a neural network happens via Gradient Descent Optimization

Linear Regression as an Optimization Problem



Objective Function

Minimize variance of
the residuals (MSE)

Linear Regression as an Optimization Problem



Objective Function

Minimize variance of
the residuals (MSE)



Constraints

Express relationship as
a straight line

$$y = Wx + b$$

Linear Regression as an Optimization Problem



Objective Function


Minimize variance of
the residuals (MSE)



Constraints

Express relationship as
a straight line


$$y = Wx + b$$




Decision Variables

Values of W and b


Minimizing MSE




Minimizing MSE




Minimizing MSE




Minimizing MSE



Minimizing MSE




Minimizing MSE




“Gradient Descent”

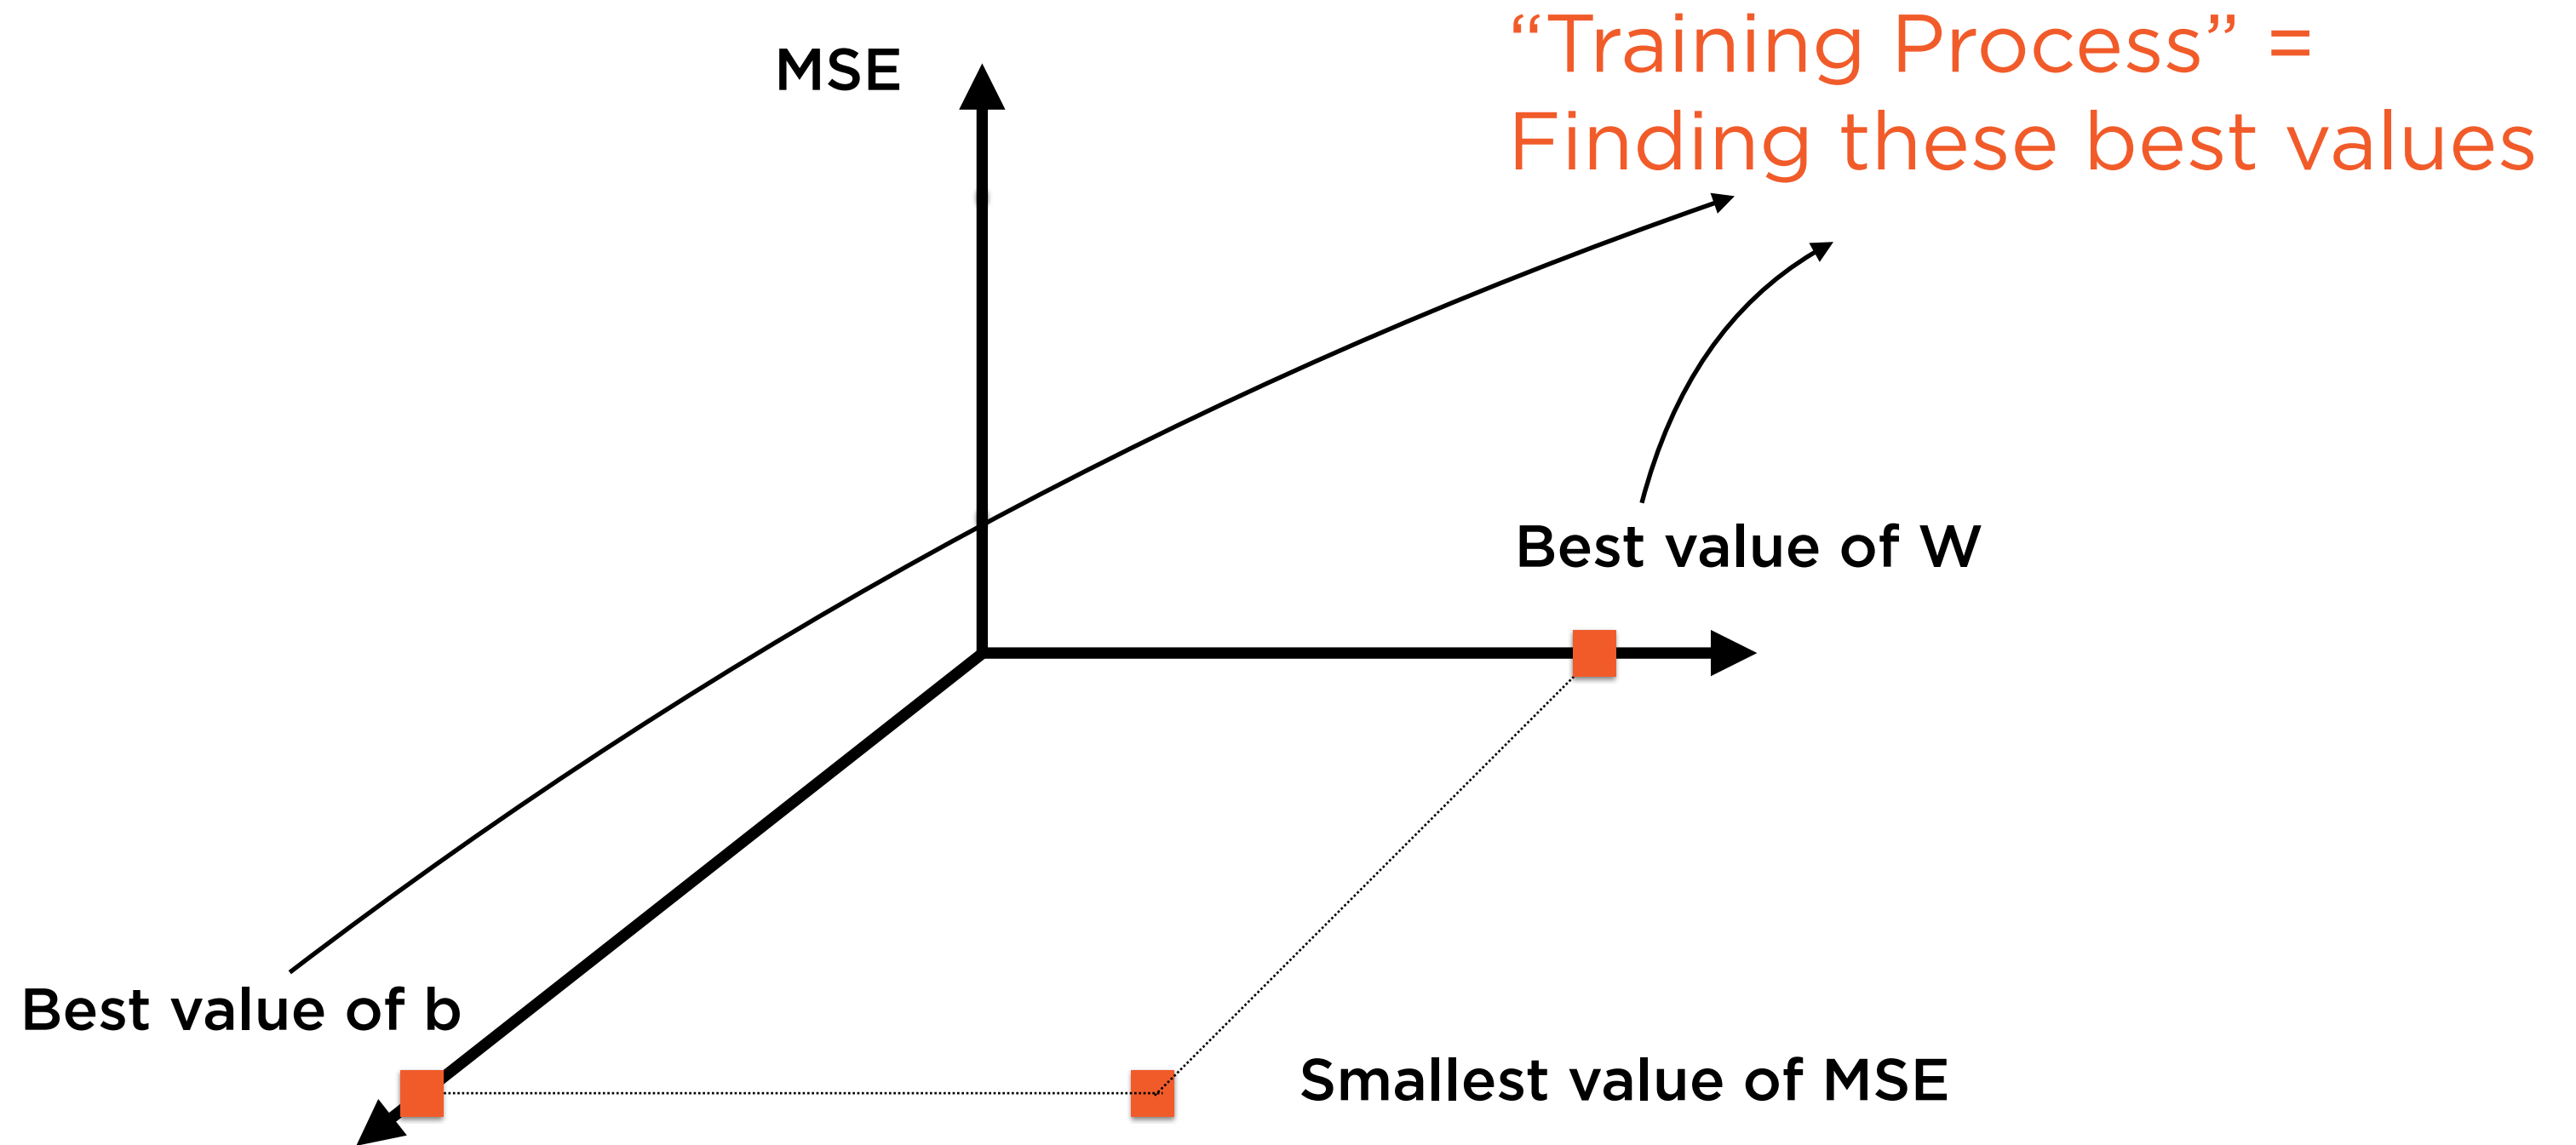
Converging on the “best”
value using an
optimization algorithm



Minimizing MSE



“Training” the Algorithm



$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{learning_rate} \times \text{Gradient}(\theta^t)$$

Basic SGD Optimizer

Move each parameter value in the direction of reducing gradient

$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{learning_rate} \times \text{Gradient}(\theta^t)$$

Basic SGD Optimizer

The size of the step in the direction of reducing gradients

$$\text{momentum_vec}^t = \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t)$$
$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{momentum_vec}$$

Momentum-based Optimizer

Momentum vector helps accelerate in the direction where gradient is decreasing

$$\begin{aligned}\text{momentum_vec}^t &= \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t) \\ \text{Parameters}^{t+1} &= \text{Parameters}^t - \text{momentum_vec}\end{aligned}$$

Momentum-based Optimizer

Gradients at each step weighted by those in previous step

Benefit: Faster convergence

$$\text{momentum_vec}^t = \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t)$$
$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{momentum_vec}$$

Momentum-based Optimizer

Need a **momentum coefficient**, between 0 and 1 to prevent overshooting

scikit-learn Optimizers

Relatively small number of optimizers supported

Plain-vanilla SGD optimizer

Adam optimizer

L-BFGS optimizer

Neural Networks in scikit-learn

**Multi-layer Perceptrons
(MLP)**

**Restricted Boltzmann
Machines (RBM)**

Neural Networks in scikit-learn

**Multi-layer Perceptrons
(MLP)**

**Restricted Boltzmann
Machines (RBM)**


MLPRegressor estimator object to build regression models

Demo


**Build and train a regression model
using neural networks in scikit-learn
using the MLPRegressor estimator**

Classification Using Neural Networks

Linear Regression with One Neuron



Linear Regression with One Neuron



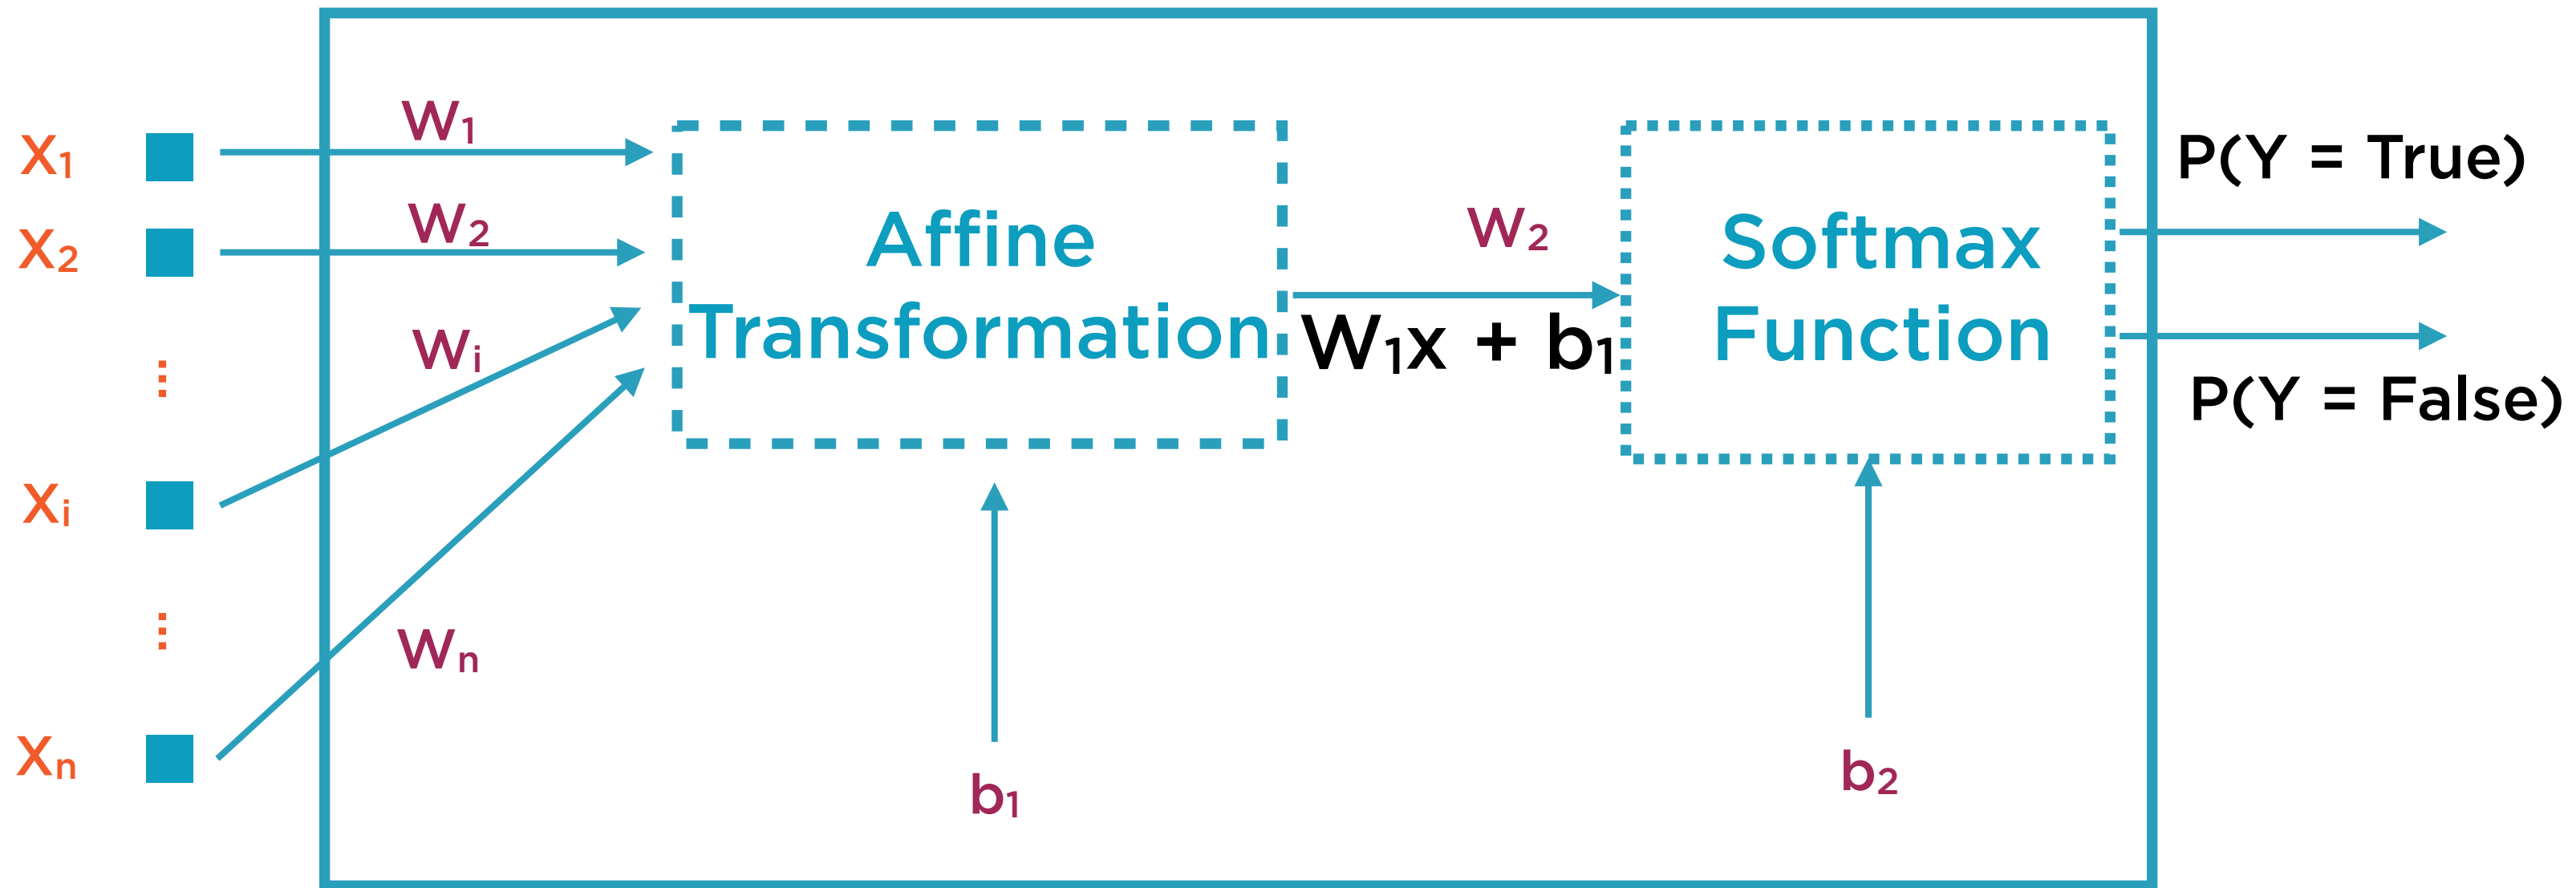
Linear Regression with One Neuron




Linear Classification with One Neuron



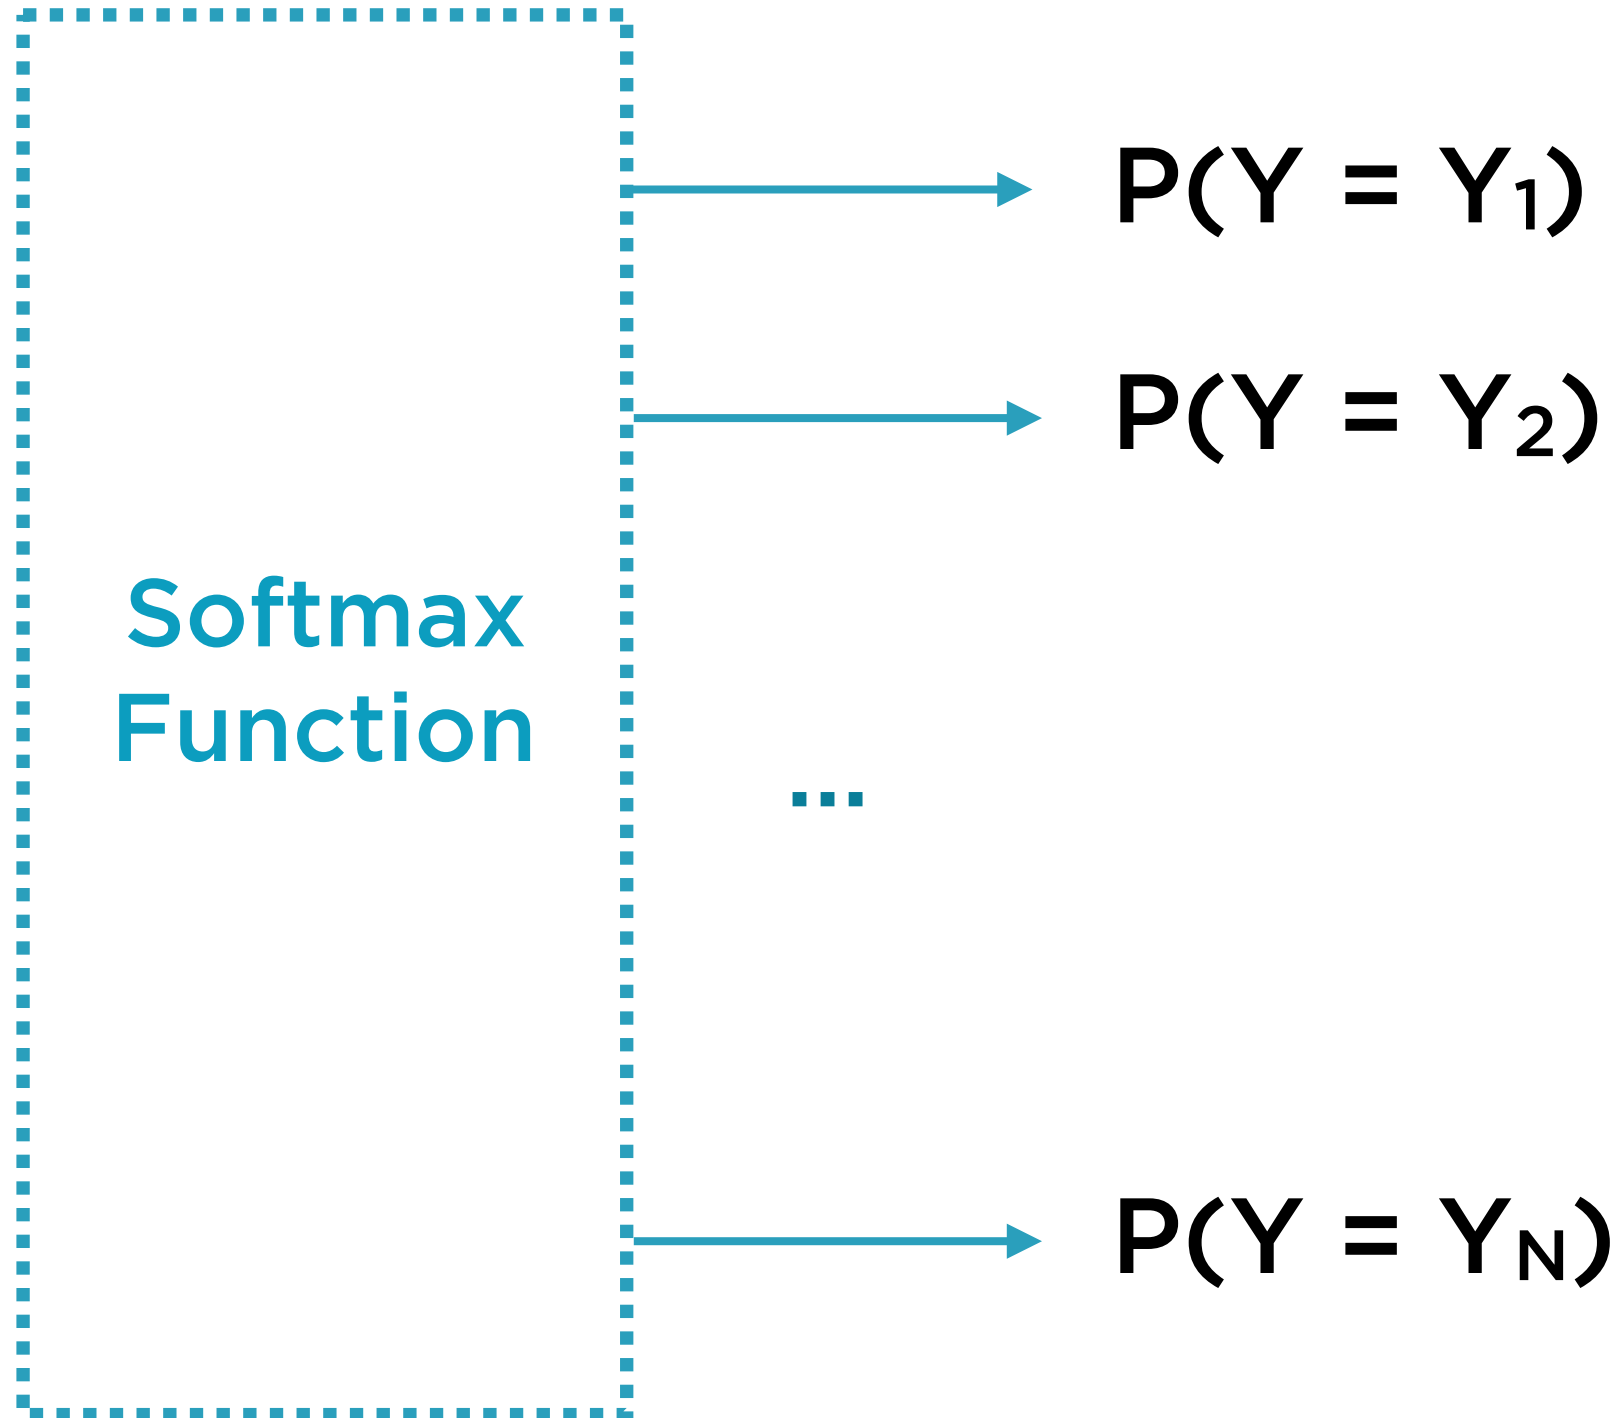
Linear Classification with One Neuron



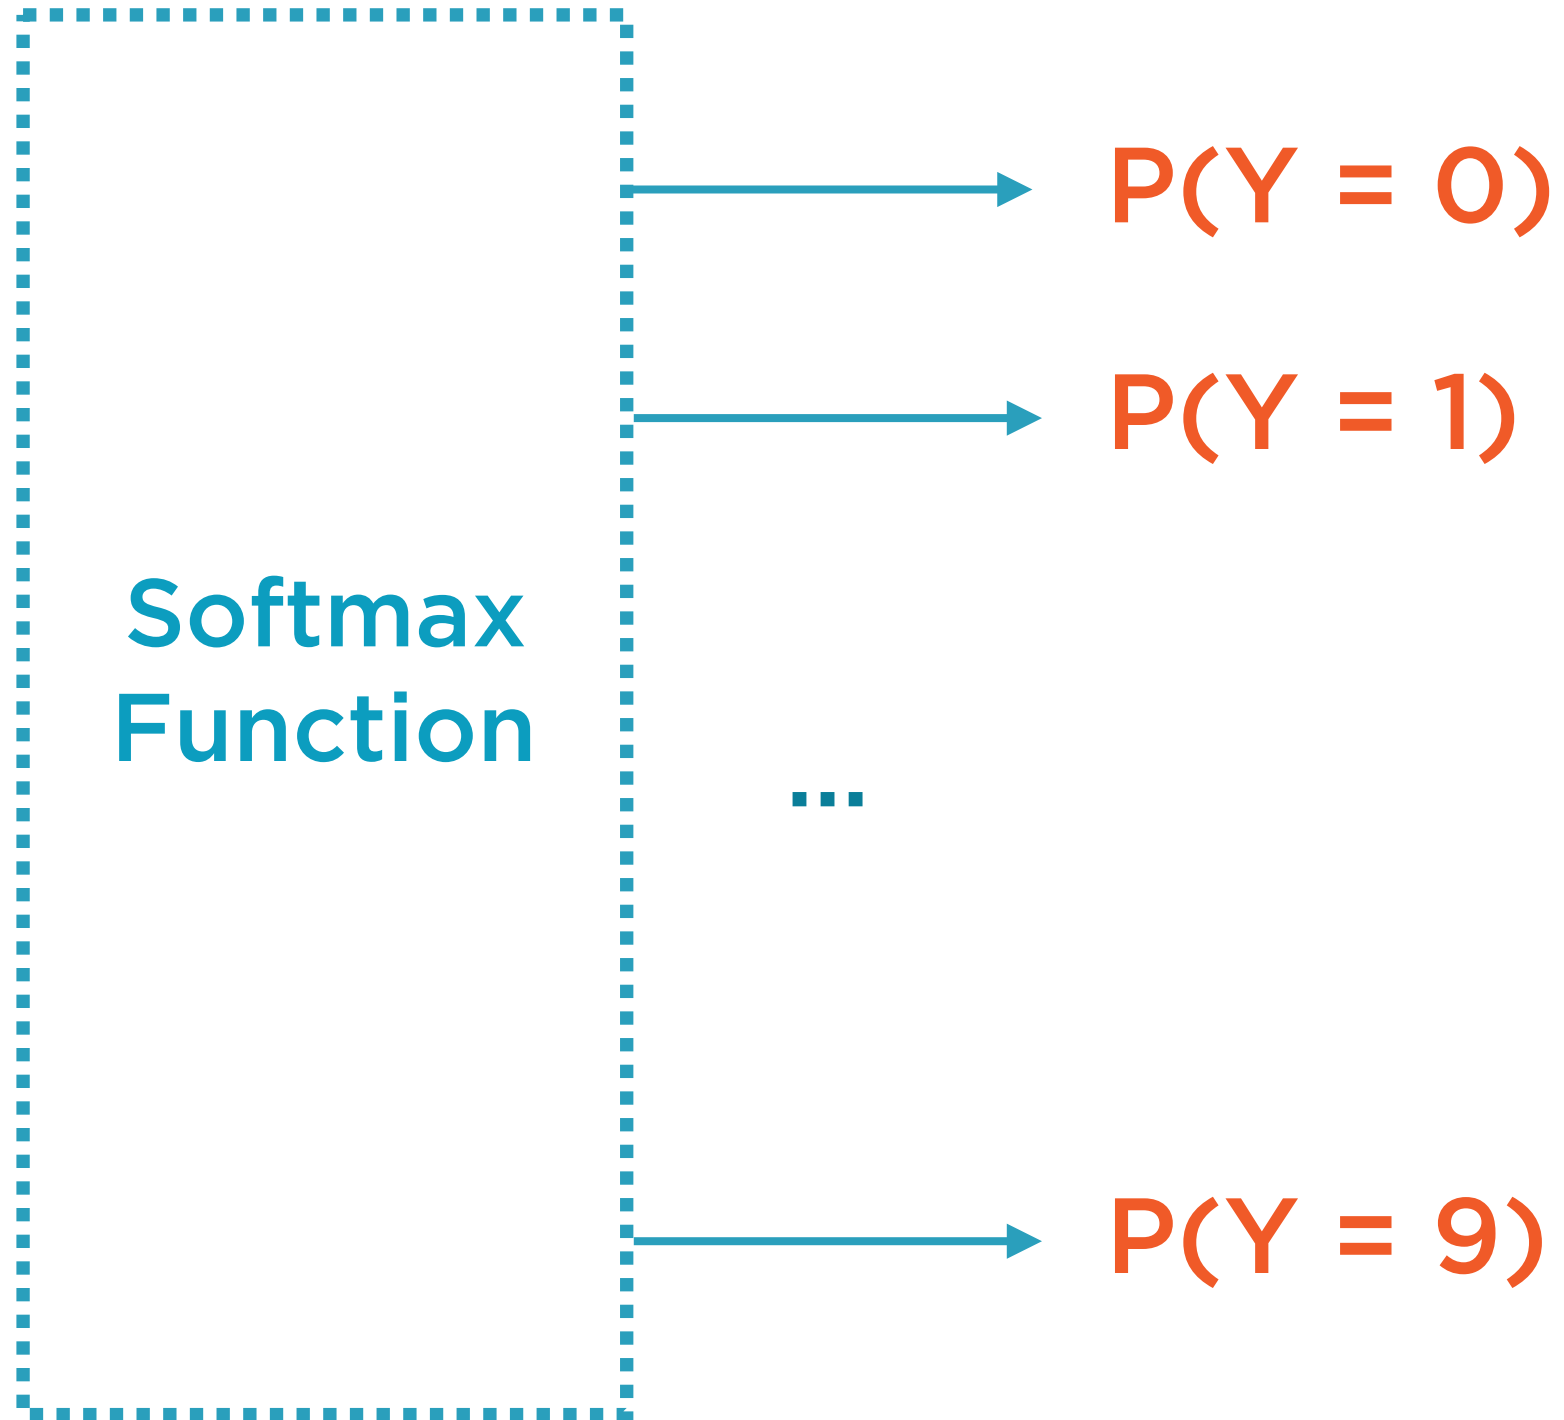
SoftMax for True/False Classification



SoftMax N-category Classification




SoftMax for Digit Classification




Intuition: Low Cross Entropy


Y_{actual}



$Y_{\text{predicted}}$




Intuition: Low Cross Entropy




The labels of the two series are in-synch

Intuition: High Cross Entropy


Y_{actual}



$Y_{\text{predicted}}$



Intuition: High Cross Entropy



The labels of the two series are out-of-synch

Demo

**Build and train a classification model
using neural networks in scikit-learn
using the MLPClassifier estimator**

Summary

Regression models using neural networks in scikit-learn

Mean square error (MSE) loss function

Using the MLPRegressor estimator to build regression models

Classification models using neural networks in scikit-learn

Cross-entropy loss function

Using the MLPClassifier estimator to build classification models